



a Gentle Introduction to Docker and All Things Containers







Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future







Devs

- all languages
- all databases
- all O/S
- targetting Linux systems

Docker will eventually be able to target FreeBSD, Solaris, and maybe OS X.







Ops

- any distro¹
- any cloud²
- any machine (physical, virtual...)
- recent kernels³



¹ as long as it's Ubuntu or Debian ⊚ others coming soon ² as long as they don't ship with their custom crappy kernel ³ at least 3.8; support for RHEL 2.6.32 on the way





CFO, CIO, CTO, ...

- LESS overhead!
- MOAR consolidation!
- MOAR agility!
- LESS costs!







Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future







The Matrix From Hell

django web frontend	?	?	?	?	?	?
node.js async API	?	?	?	?	?	?
background workers	?	?	?	?	?	?
SQL database	?	?	?	?	?	?
distributed DB, big data	?	?	?	?	?	?
message queue	?	?	?	?	?	?
	my laptop	your laptop	QA	staging	prod on cloud VM	prod on bare metal







Another Matrix from Hell



?

?

?

?

?

?

?



?

?

?

?

?

?

?



?

?

?

?

?

?

?



?

?

?

?

?

?

7



?

?

?

?

?

?

?



?

7

?

2

2

?

2



















Solution: the *intermodal shipping container*

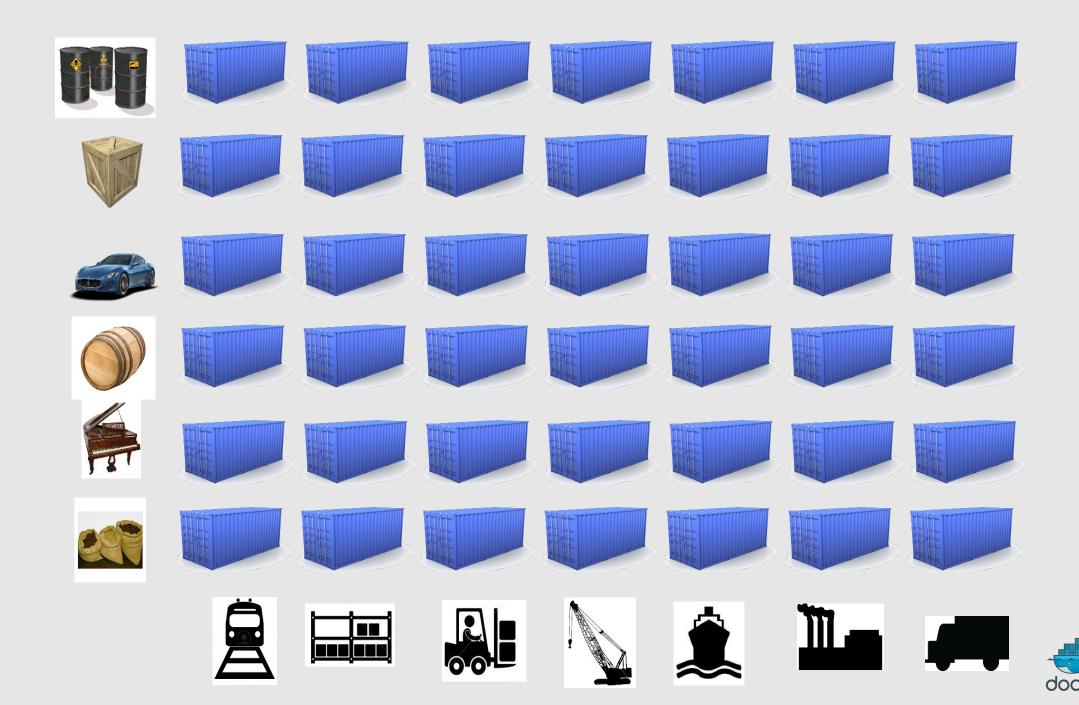








Solved!

















Linux containers...

Units of software delivery (ship it!)

- run everywhere
 - regardless of kernel version
 - regardless of host distro
 - (but container and host architecture must match*)
- run anything
 - if it can run on the host, it can run in the container
 - i.e., if it can run on a Linux kernel, it can run



^{*}Unless you emulate CPU with qemu and binfmt



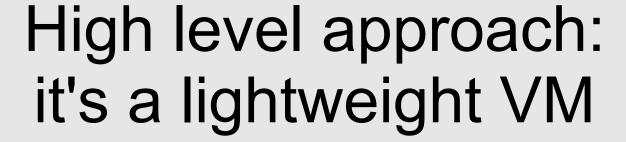


Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future







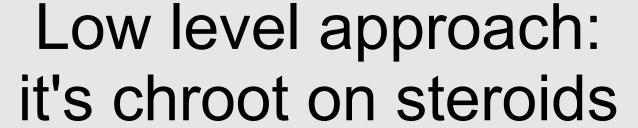


- own process space
- own network interface
- can run stuff as root
- can have its own /sbin/init (different from the host)

« Machine Container »









- can also not have its own /sbin/init
- container = isolated process(es)
- share kernel with host
- no device emulation (neither HVM nor PV)

« Application Container »





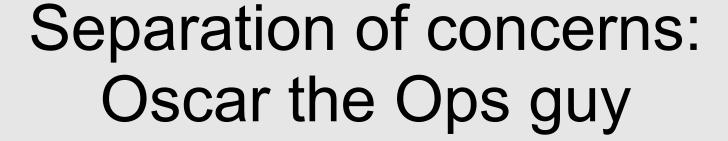




- inside my container:
 - my code
 - my libraries
 - my package manager
 - my app
 - my data





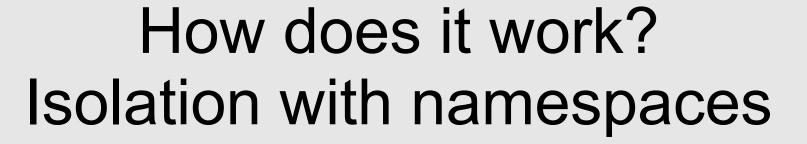




- outside the container:
 - logging
 - remote access
 - network configuration
 - monitoring





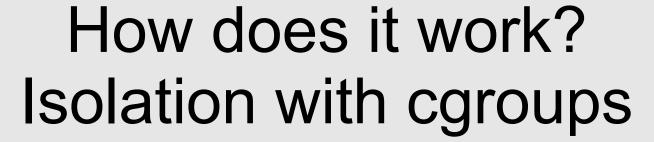




- pid
- mnt
- net
- uts
- ipc
- user





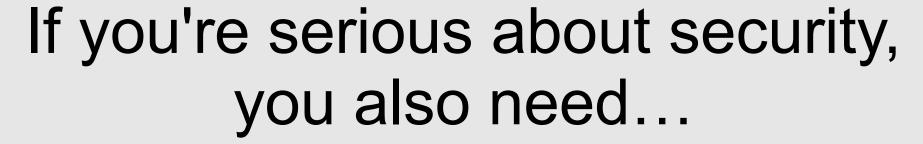




- memory
- cpu
- blkio
- devices





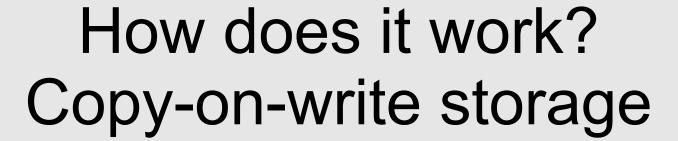




- capabilities
 - okay: cap_ipc_lock, cap_lease, cap_mknod, cap_net_admin, cap_net_bind_service, cap_net_raw
 - troublesome: cap_sys_admin (mount!)
- think twice before granting root
- grsec is nice
- seccomp (very specific use cases); seccomp-bpf
- beware of full-scale kernel exploits!









- unioning filesystems (AUFS, overlayfs)
- snapshotting filesystems (BTRFS, ZFS)
- copy-on-write block devices (thin snapshots with LVM or device-mapper)

This is now being integrated with low-level LXC tools as well!



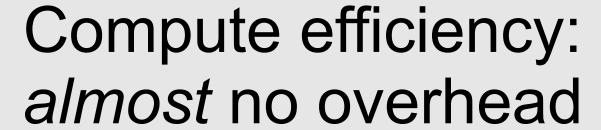




Efficiency









- processes are isolated, but run straight on the host
- CPU performance= native performance
- memory performance
 a few % shaved off for (optional) accounting
- network performance
 - = small overhead; can be reduced to zero







Storage efficiency: many options!

	Union Filesystems	Snapshotting Filesystems	Copy-on-write block devices
Provisioning	Superfast Supercheap	Fast Cheap	Fast Cheap
Changing small files	Superfast Supercheap	Fast Cheap	Fast Costly
Changing large files	Slow (first time) Inefficient (copy-up!)	Fast Cheap	Fast Cheap
Diffing	Superfast	Superfast (ZFS) Kinda meh (BTRFS)	Slow
Memory usage	Efficient	Efficient	Inefficient (at high densities)
Drawbacks	Random quirks AUFS not mainline !AUFS more quirks	ZFS not mainline BTRFS not as nice	Higher disk usage Great performance (except diffing)
Bottom line	Ideal for PAAS and high density things	This might be the Future	Dodge Ram 3500







Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future















Docker-what?

- Open Source engine to commoditize LXC
- using copy-on-write for quick provisioning

STOP! HAMER DEMOTIME.









Yes, but...

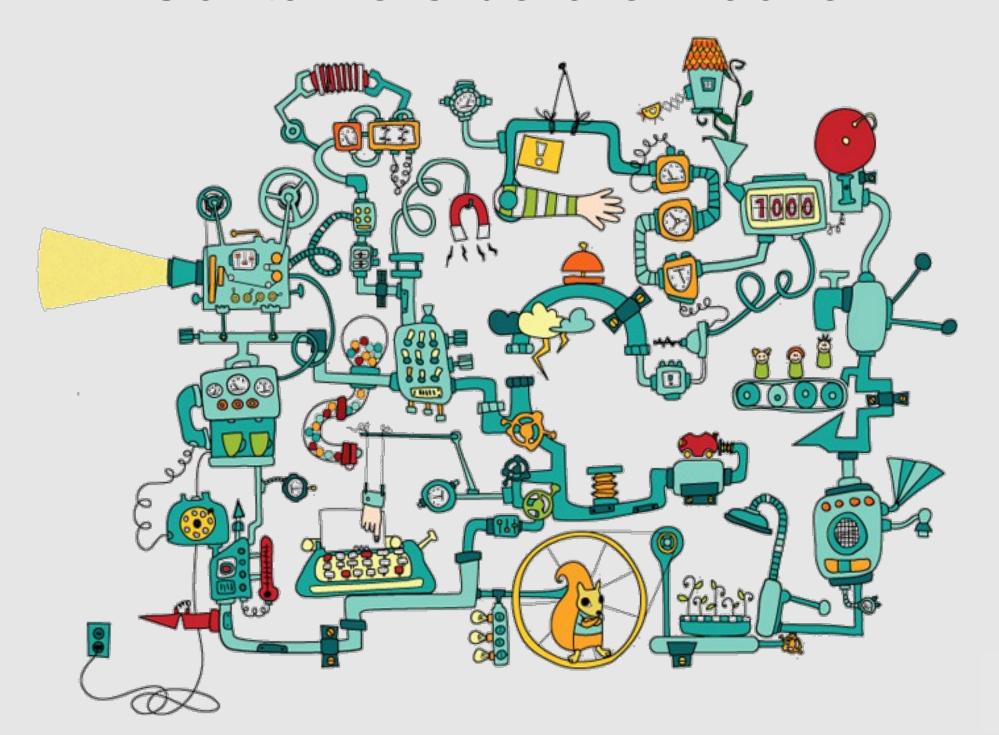
- « I don't need Docker;
 I can do all that stuff with LXC tools, rsync, some scripts! »
- correct on all accounts;
 but it's also true for apt, dpkg, rpm, yum, etc.
- the whole point is to commoditize,
 i.e. make it ridiculously easy to use







Containers before Docker

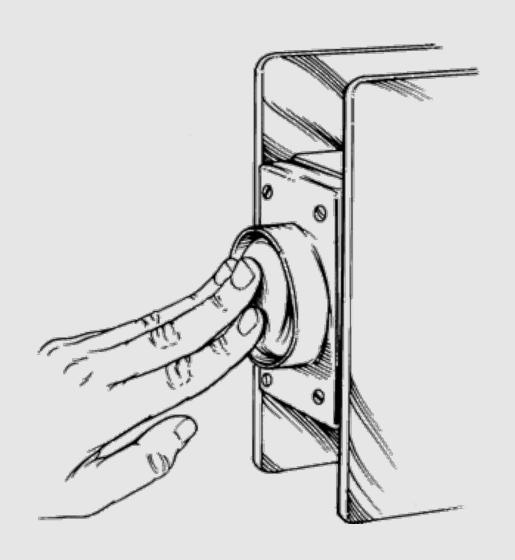








Containers after Docker









What this really means...

- instead of writing « very small shell scripts » to manage containers, write them to do the rest:
 - continuous deployment/integration/testing
 - orchestration
- = use Docker as a building block
- re-use other people images (yay ecosystem!)









- Open Source engine to commoditize LXC
- using copy-on-write for quick provisioning
- allowing to create and share images
- standard format for containers
 (stack of layers; 1 layer = tarball+metadata)
- standard, reproducible way to easily build trusted images (Dockerfile, Stackbrew...)





Docker-what? Under The Hood



- rewrite of dotCloud internal container engine
 - original version: Python, tied to dotCloud's internal stuff
 - released version: Go, legacy-free
- the Docker daemon runs in the background
 - manages containers, images, and builds
 - HTTP API (over UNIX or TCP socket)
 - embedded CLI talking to the API
- Open Source (GitHub public repository + issue tracking)
- user and dev mailing lists
- FreeNode IRC channels #docker, #docker-dev







Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future





Authoring images with run/commit

- 1) docker run ubuntu bash
- 2) apt-get install this and that
- 3) docker commit <containerid> <imagename>
- 4) docker run <imagename> bash
- 5) git clone git://.../mycode
- 6) pip install -r requirements.txt
- 7) docker commit <containerid> <imagename>
- 8) repeat steps 4-7 as necessary
- 9) docker tag <imagename> <user/image>
- 10) docker push <user/image>







Authoring images with a Dockerfile

FROM ubuntu

docker build -t jpetazzo/couchdb .









0) create a GitHub account

On index.docker.io:

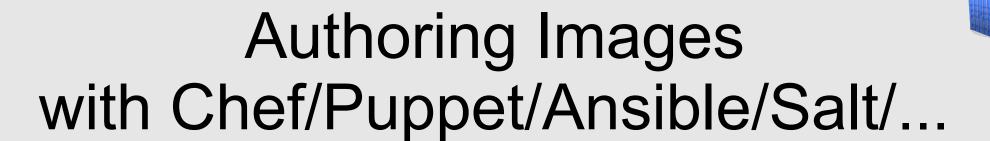
- 1) create a Docker account
- 2) link it with your GitHub account
- 3) enable Trusted Builds on any public repo

On your dev env:

- 4) git add Dockerfile
- 5) git commit
- 6) git push







Plan A: « my other VM is a container »

- write a Dockerfile to install \$YOUR_CM
- start tons of containers
- run \$YOUR_CM in them

Good if you want a mix of containers/VM/metal But slower to deploy, and uses more resources







Plan B: « the revolution will be containerized »

- write a Dockerfile to install \$YOUR_CM
- ... and run \$YOUR_CM as part of build process
- deploy fully baked images

Faster to deploy

Easier to rollback







Outline

- Whom is this for?
- What's the problem?
- What's a Container?
- Docker 101
- Docker images
- Docker deployment
- Docker future







Running containers

- SSH to Docker host and manual pull+run
- REST API (feel free to add SSL certs, OAUth...)
- OpenStack Nova
- OpenStack Heat
- who's next? OpenShift, CloudFoundry?
- multiple Open Source PAAS built on Docker (Cocaine, Deis, Flynn...)





Orchestration & Service Discovery (0.6.5)

- you can name your containers
- they get a generated name by default (red_ant, gold_monkey...)
- you can link your containers

```
docker run -d -name frontdb
docker run -d -link frontdb:sql frontweb
```

→ container frontweb gets one bazillion environment vars





Orchestration & Service Discovery roadmap

- currently single-host
- problem: how do I link with containers on other hosts?
- solution: ambassador pattern!
 - app container runs in its happy place
 - other things (Docker, containers...) plumb it







Orchestration roadmap

- currently static
- problem: what if I want to...
 move a container?
 do a master/slave failover?
 WebScale my MangoDB cluster?
- solution: dynamic discovery!

