Kubernetes in Docker

Alex Mavrogiannis

Docker EE Engineering

Agenda

1. Introduction
2. Demo: Kubernetes in Docker EE 2.0
3. General CE/EE Architectures
4. EE: Topics on mixed workloads
5. EE: AuthN/AuthZ
6. Q&A

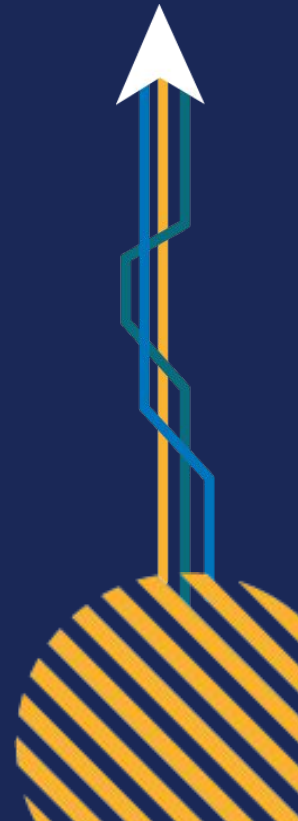

Introduction

What are Docker containers?

Processes running on the same host OS using the following mechanisms:

- IPC Namespaces
- PID Namespaces
- Network Namespaces
- Control Groups (Memory/CPU)
- Union Filesystems and image distribution mechanisms

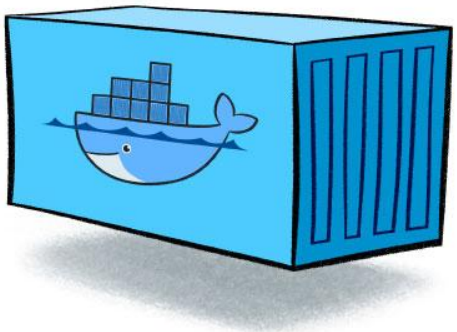

Containers are managed by the Container Runtime process running on the host OS, the Docker Engine.

What is a container orchestrator?

Management of containers running in one or more container runtimes

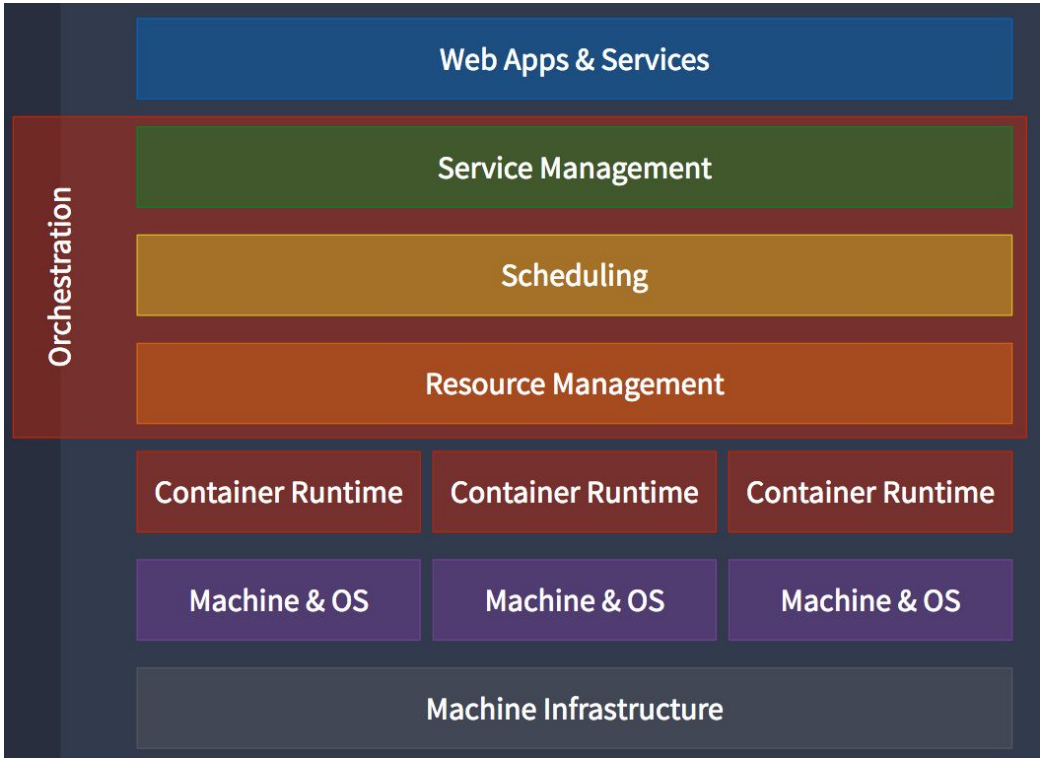

Orchestrator: Docker Swarm

- github.com/docker/swarm
- Cluster-wide imperative API based on the single-node API of the Docker Engine
- High Availability and peer discovery managed through a pluggable discovery backend: etcd, consul
- “Leader” caches entire cluster state: containers, volumes, networks etc.

Orchestrator: Docker Engine with Swarm-Mode Enabled

- github.com/docker/swarmkit
- Declarative State through the “Service” construct
- Built-in Routing Mesh & Overlay networking
- In-memory Raft Store for all state (persisted to disk)
- Built-in CA, per-node cryptographic node identity, mTLS between all endpoints

Orchestrator: Kubernetes

- github.com/kubernetes/kubernetes
- Scheduling Unit: Pods
- Declarative State through “Controllers”: Deployment, ReplicaSet, DaemonSet …
- Flat Networking model delegated to plugins

Docker: Now Powered by Swarm and Kubernetes

Docker EE 2.0: A conformant kubernetes distribution

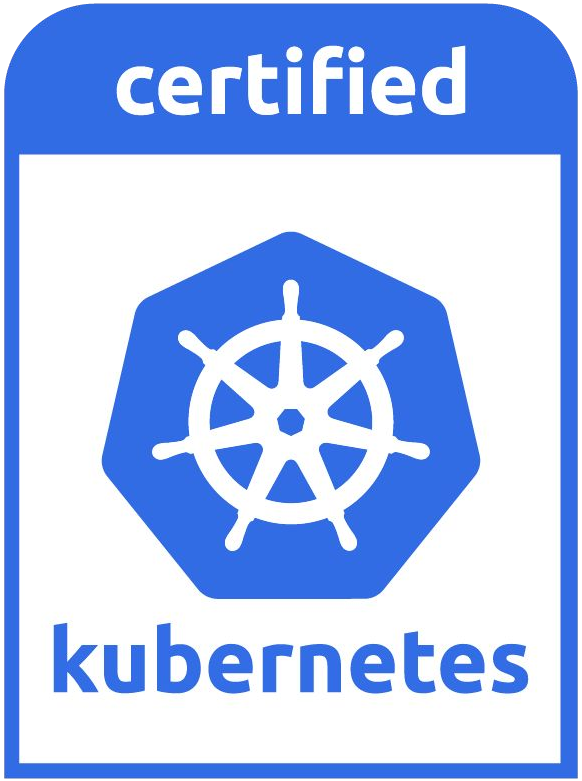

Demo: Kubernetes in Docker EE 2.0

General CE/EE Architecture

Kubernetes in Docker CE (Windows and Mac)

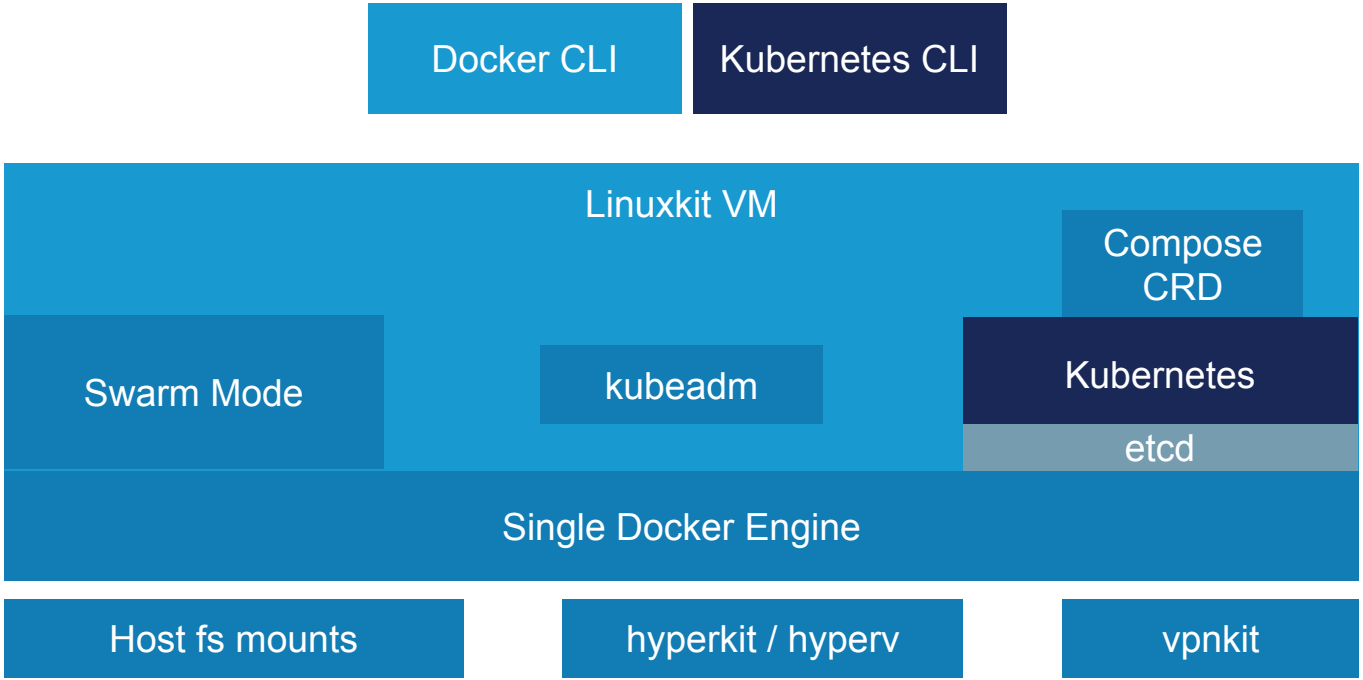

Docker EE to include Kubernetes

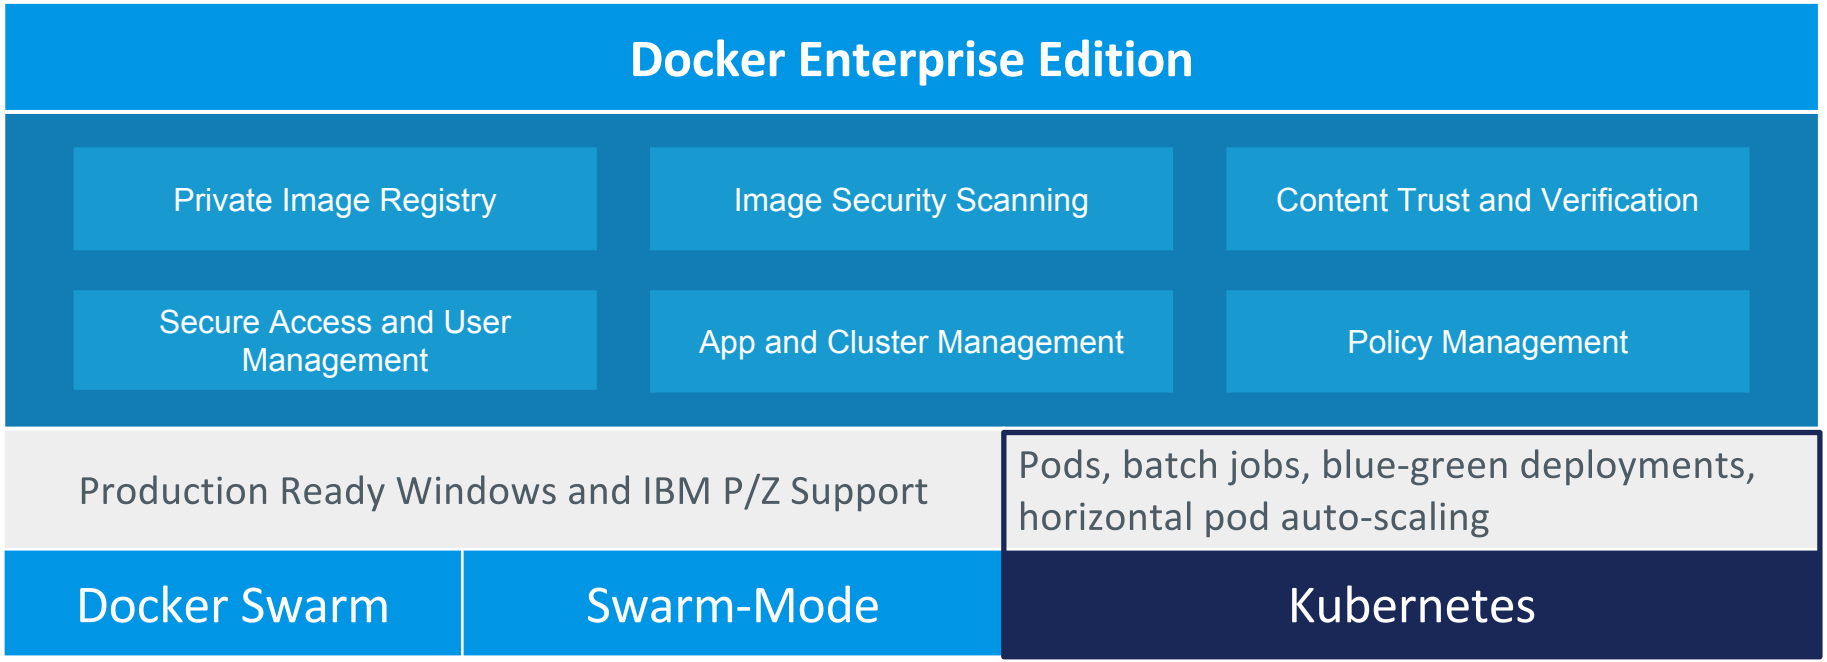

Kubernetes in Docker EE

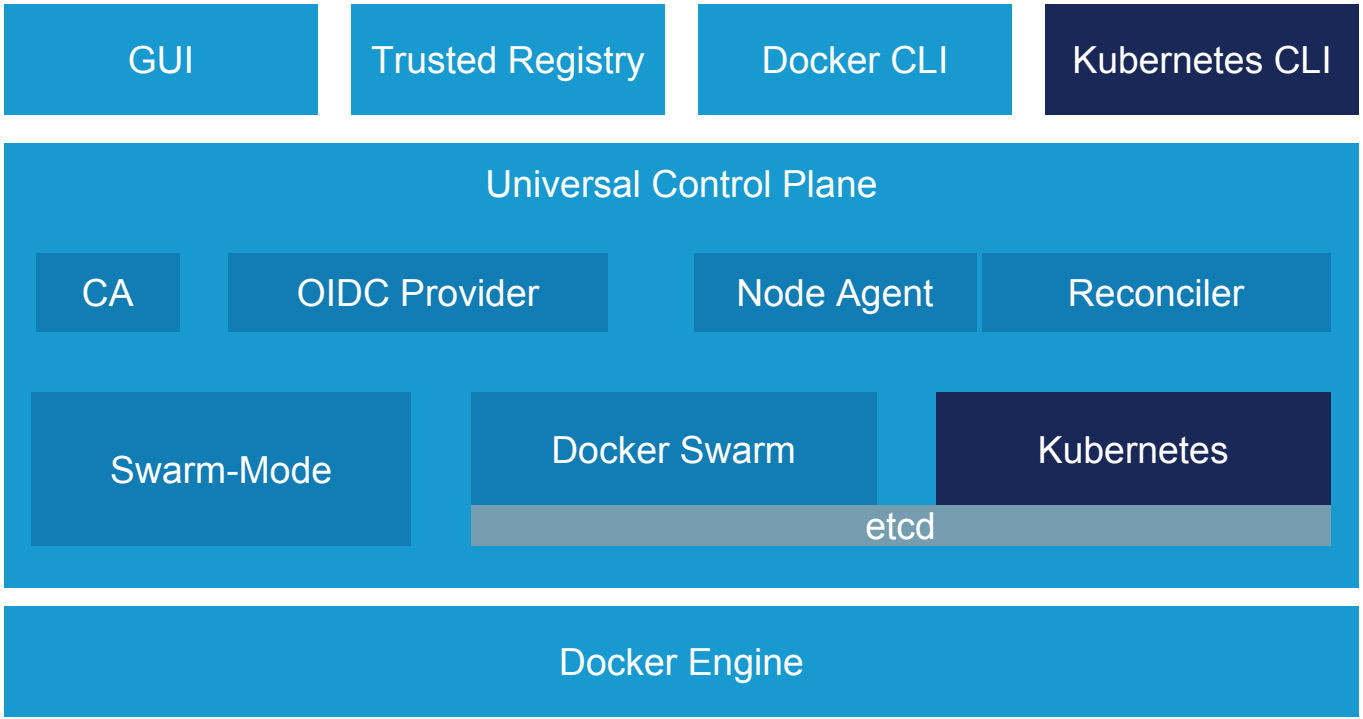

Docker EE Architectural Highlights

- Conformant Kubernetes components ran as Docker containers
- Swarm Managers are Kubernetes Masters
- Swarmkit node inventory is source of truth
- Cryptographic Node Identity and mTLS used throughout

Kubernetes Plugin Interfaces in Docker EE

- General:
  - Native API extensibility supported
  - Some apiserver/kubelet flags modifiable by users
- Networking:
  - Support for CNI plugin during install
  - Ingress
- Storage: Docker Volume Plugins supported via built-in flexvolume driver, CSI in future
- Metrics: Heapster Storage Backends or Prometheus

Topics on Mixed Workloads

Resource Contention

- Allocatable Resources: The set of CPU and Memory resources available for scheduling by an orchestrator on a single node
- Multiple orchestrators = Different definitions of allocatable resources
    - Docker Swarm: Respectful of CPU/Memory limits, but container cache may be stale
    - Docker Engine with Swarm-Mode: Only aware of its own reservations
    - Kubernetes: Effective handling of out-of-resource situations, but only for kubernetes workloads
- When a node is at/near capacity:
    - All CPU shares throttled equally
    - The OS's OOM killer kills processes
    - All orchestrators will reschedule on OOM, but potential workload interruption

Orchestrator Selection

- Each node is running both kubernetes and swarm system components
- Administrators can toggle between (kubernetes, swarm or mixed) scheduling for any given node.
- When toggling orchestrators, workloads of the previous orchestrator will be evicted
- If a node is not enabled for a given orchestrator, users will not be able to schedule workloads on that node using that orchestrator.

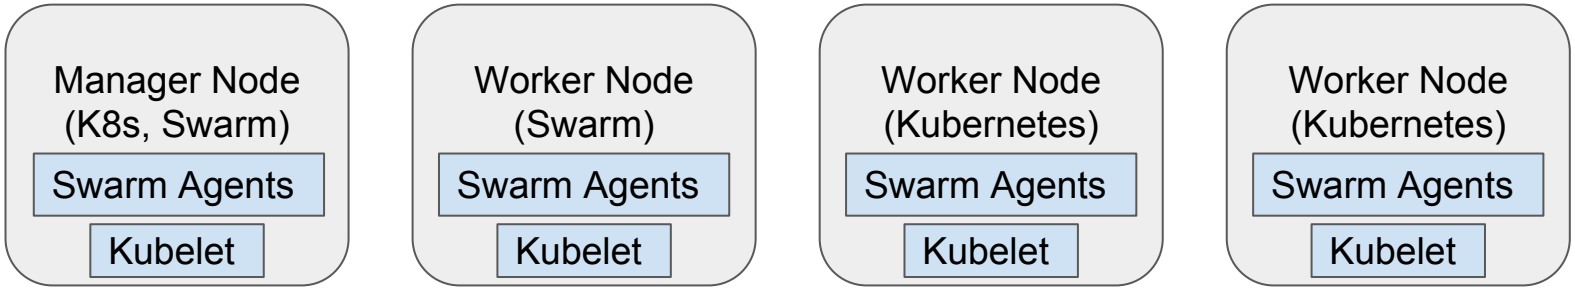

Workload Interoperability

- Networking
  - Layer 3 not connected between kubernetes & swarm
  - Batteries-included kubernetes ingress controller
  - Layer 7 routing for swarm workloads
  - Configure external DNS
- Storage: Kubernetes workloads with docker volumes via flexvolume

AuthN / AuthZ

Identity Providers

Systems that manage identity information for principals and provides user authentication as a service.

- SAML
- OpenID Connect (OIDC)

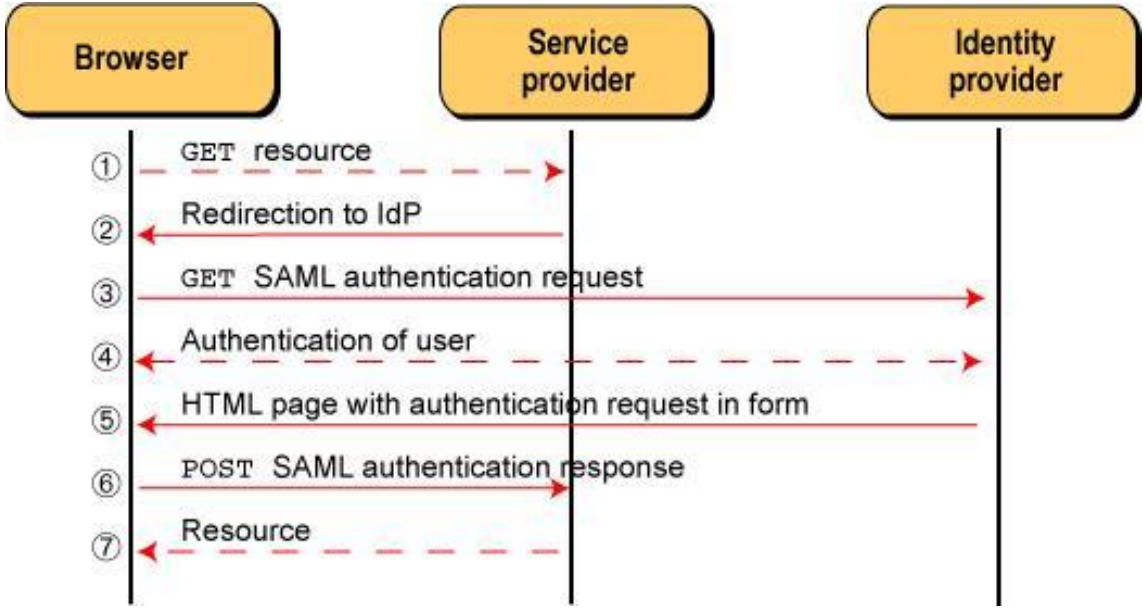

Actors in Docker EE Authentication/Authorization

- Client (Browser, Docker CLI or kubectl)
- UCP Controller
- OIDC Provider
- Kubernetes API server

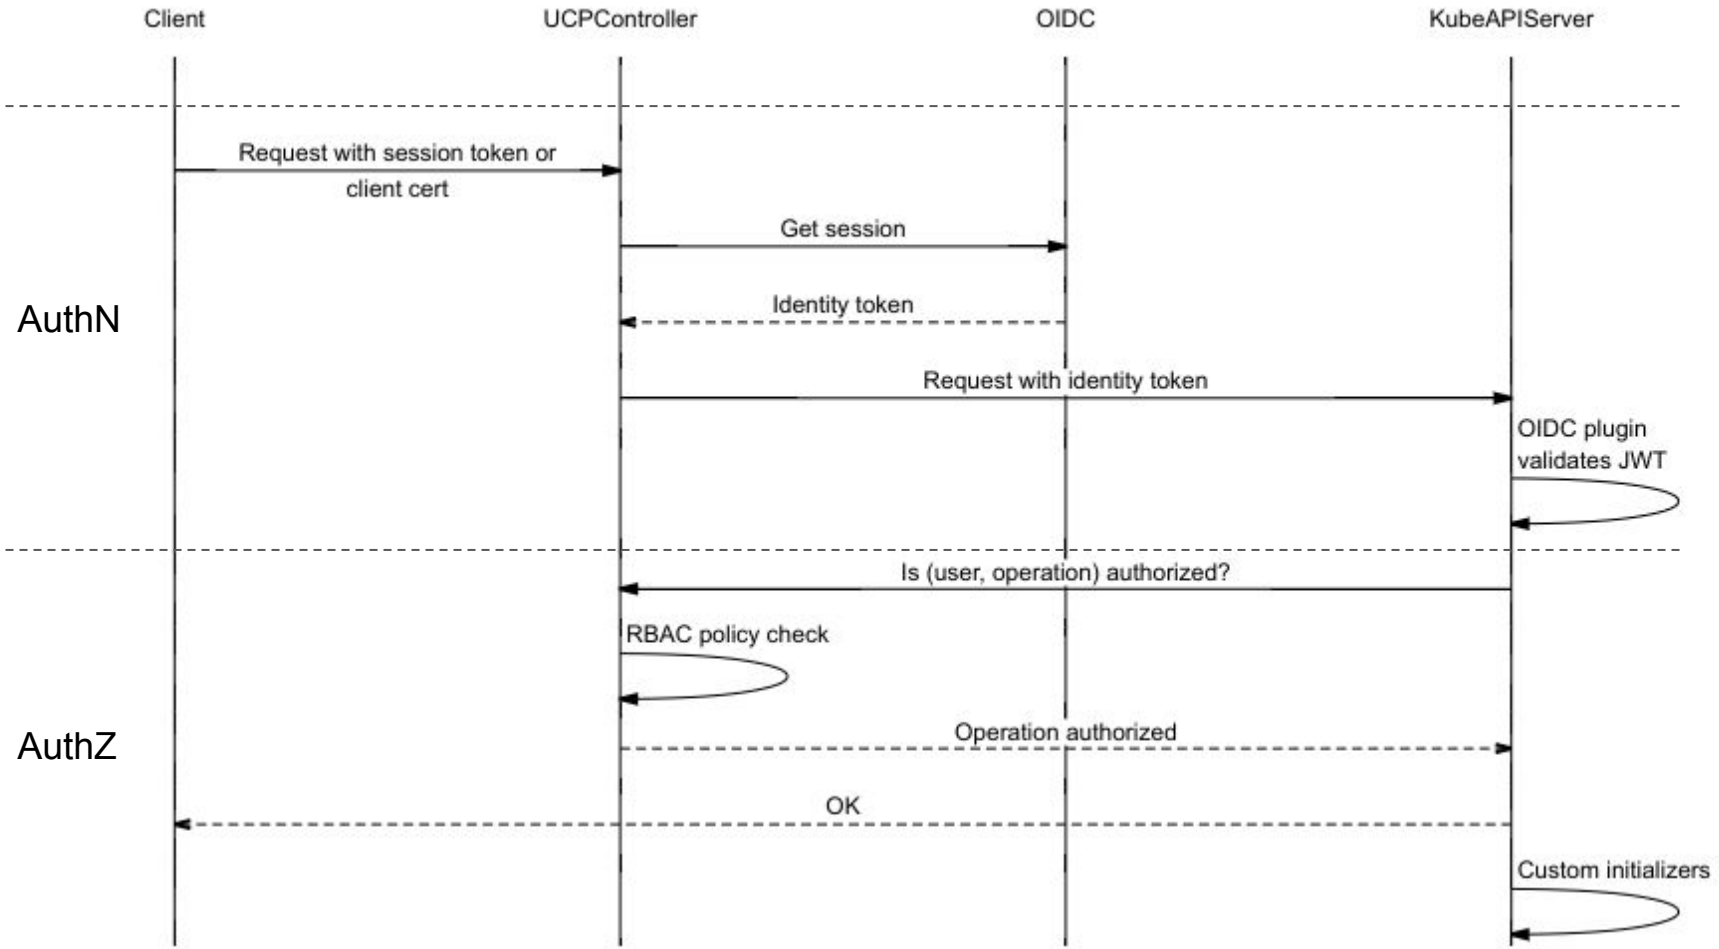
Client
UCPController
OIDC
KubeAPIServer
AuthN
Request with session token or client cert
Get session
Identity token
Request with identity token
OIDC plugin validates JWT
AuthZ
Is (user, operation) authorized?
RBAC policy check
Operation authorized
OK
Custom initializers

In Summary...

- Docker EE and CE will include a conformant Kubernetes distribution.
- Resource Contention mitigated via orchestrator selection
- In EE, Authentication and Authorization integrated via standard plugin interfaces.

Thank You!

Sign up for the beta at docker.com/kubernetes

alexmavr

alex.mavrogiannis@docker.com

dockercon17 EU