

# **Experienced**

Device driver(windows, linux)  
Media streaming  
CDN  
Docker

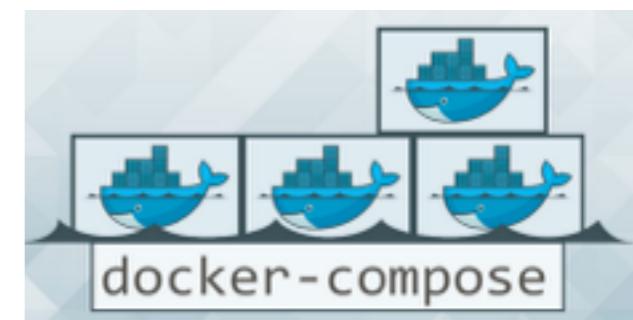
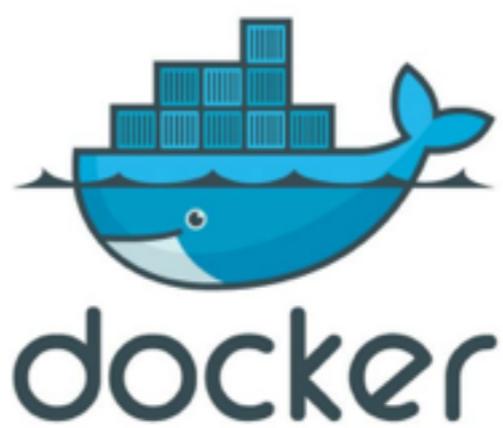
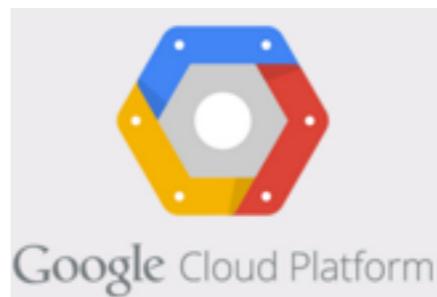
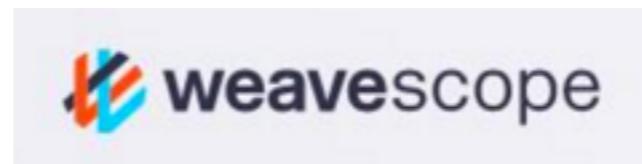
PyCon Korea 2015 Speaker  
PyCon Hongkong 2015 Speaker



**IP[y]:**  
IPython



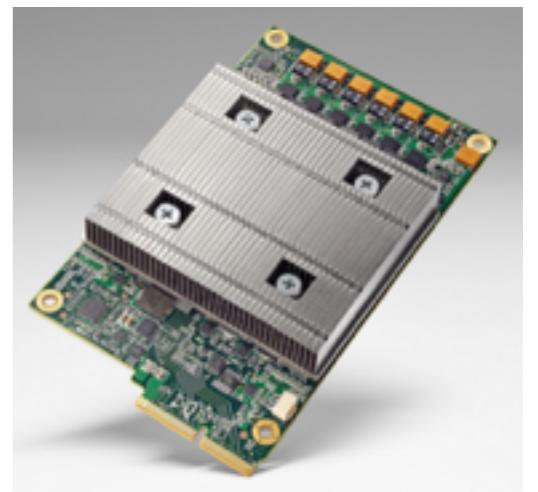
**yamL**



# Agenda

- TensorFlow + GPU
- TensorFlow Distributed Environment
- Docker, Google Cloud Platform, AWS..
- Docker, Docker-compose...
- Demo

# TensorFlow + GPU



TPU

# Different way to install TensorFlow

Pip install

Virtualenv install

Anaconda install

**Docker install**

Installing from sources

```
// Cool~  
$ docker run -it -p 8888:8888 gcr.io/tensorflow/tensorflow  
  
// GPU?  
$ path/to/repo/tensorflow/tools/docker/docker_run_gpu.sh  
gcr.io/tensorflow/tensorflow:gpu
```



<https://github.com/tensorflow/tensorflow>,

## **Branch:master,0.80,0.90rc (Binary download)**

- Linux CPU only: Python 2 (build history) / Python 3.4 (build history) / Python 3.5 (build history)
- **Linux GPU**: Python 2 (build history) / Python 3.4 (build history) / Python 3.5 (build history)
- **Mac CPU only**: Python 2 (build history) / Python 3 (build history)
- Android (build history)

// 2) install & setup for gcc and clang package

```
$ xcode-select --install
```

```
$ vi .zshrc
#####
# CUDA
export PATH=/Developer/NVIDIA/CUDA-7.5/bin:$PATH
export DYLD_LIBRARY_PATH=/Developer/NVIDIA/CUDA-7.5/lib:$DYLD_LIBRARY_PATH
```

```
$ kextstat | grep -i cuda
196 0xfffffff7f82df8000 0x2000 0x2000 com.nvidia.CUDA (1.1.0) 5AFE550D-6361-3897-912D-897C13FF6983 <4 1>
```

```
// nvcc compiler
$ nvcc -V
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2015 NVIDIA Corporation
Built on Mon_Apr_11_13:23:40_CDT_2016
Cuda compilation tools, release 7.5, V7.5.26
```

```
$ cd /Developer/NVIDIA/CUDA-7.5/samples/
```

```
$ sudo make -C 0_Simple/vectorAdd
$ sudo make -C 0_Simple/vectorAddDrv
$ sudo make -C 1_Utils/deviceQuery
$ sudo make -C 1_Utils/bandwidthTest
```

```
$ cd bin/x86_64/darwin/release
$ ./deviceQuery
./deviceQuery
./deviceQuery Starting...
```

#### CUDA Device Query (Runtime API) version (CUDART static linking)

```
cudaGetDeviceCount returned 35
-> CUDA driver version is insufficient for CUDA runtime version
Result = FAIL
```

```
$ ./bandwidthTest
CUDA Bandwidth Test] - Starting...
Running on...
```

```
cudaGetDeviceProperties returned 35
-> CUDA driver version is insufficient for CUDA runtime version
CUDA error at bandwidthTest.cu:255 code=35(cudaErrorInsufficientDriver) "cudaSetDevice(currentDevice)"
```

1) download & install  
GPU support : CUDA

<https://developer.nvidia.com/cuda-downloads>

This Mac  
OS X El Capitan  
Version 10.11.4  
MacBook Pro ( Retina, 15-inch, Mid 2015)  
Processor 2.2 GHz Intel Core i7  
Memory 16GB 1600 MHz DDR3  
Graphics Intel Iris Pro 1536 MB

[Code](#) [Issues 363](#) [Pull requests 28](#)[Filters](#) [GPU is:open](#)[Clear current search query, filters, and sorts](#)[124 Open](#) ✓ 496 Closed[GPU Host To Device copies don't appear to](#)  
#2848 opened an hour ago by sjperkins[segfault in perftools::gputools::StreamExecutor::DeviceMemoryUsage - on a busy gpu](#)  
#2840 opened 15 hours ago by Mlinner[Can not completely disable GPU & CUDA support.](#)  
#2829 opened 21 hours ago by MrSlayer02[ERROR when building pip package: undeclared function related to swig](#)  
#2811 opened 2 days ago by peterzcc[CUDA\\_ERROR\\_MISALIGNED\\_ADDRESS on MNIST example](#)  
#2810 opened 3 days ago by johnfrombluff[Slow quantized graph](#)  
#2807 opened 3 days ago by changyun79[Building TF with custom GCC requires hardcoded ld,nm and as](#)  
#2806 opened 3 days ago by akors[error when using batch\\_normalize in tensorflow.contrib](#)  
#2797 opened 3 days ago by guohengkai[Tensorflow hangs without explanation, some threads stay busy.](#) awaiting response  
#2788 opened 4 days ago by alexatkin1[Potential Contribution: Ansible roles & playbooks](#)  
#2780 opened 4 days ago by kuza55[TensorFlow : Machine Translation example Segmentation fault \(Core dumped\)](#) awaiting response  
#2776 opened 4 days ago by suraj1990[Segmentation fault on tensorflow 0.9.0](#) awaiting response  
#2773 opened 4 days ago by amineHorseman[output ImportError libcudart.so.7.5. in window 7 pycharm](#) awaiting response  
#2771 opened 4 days ago by ismymajia

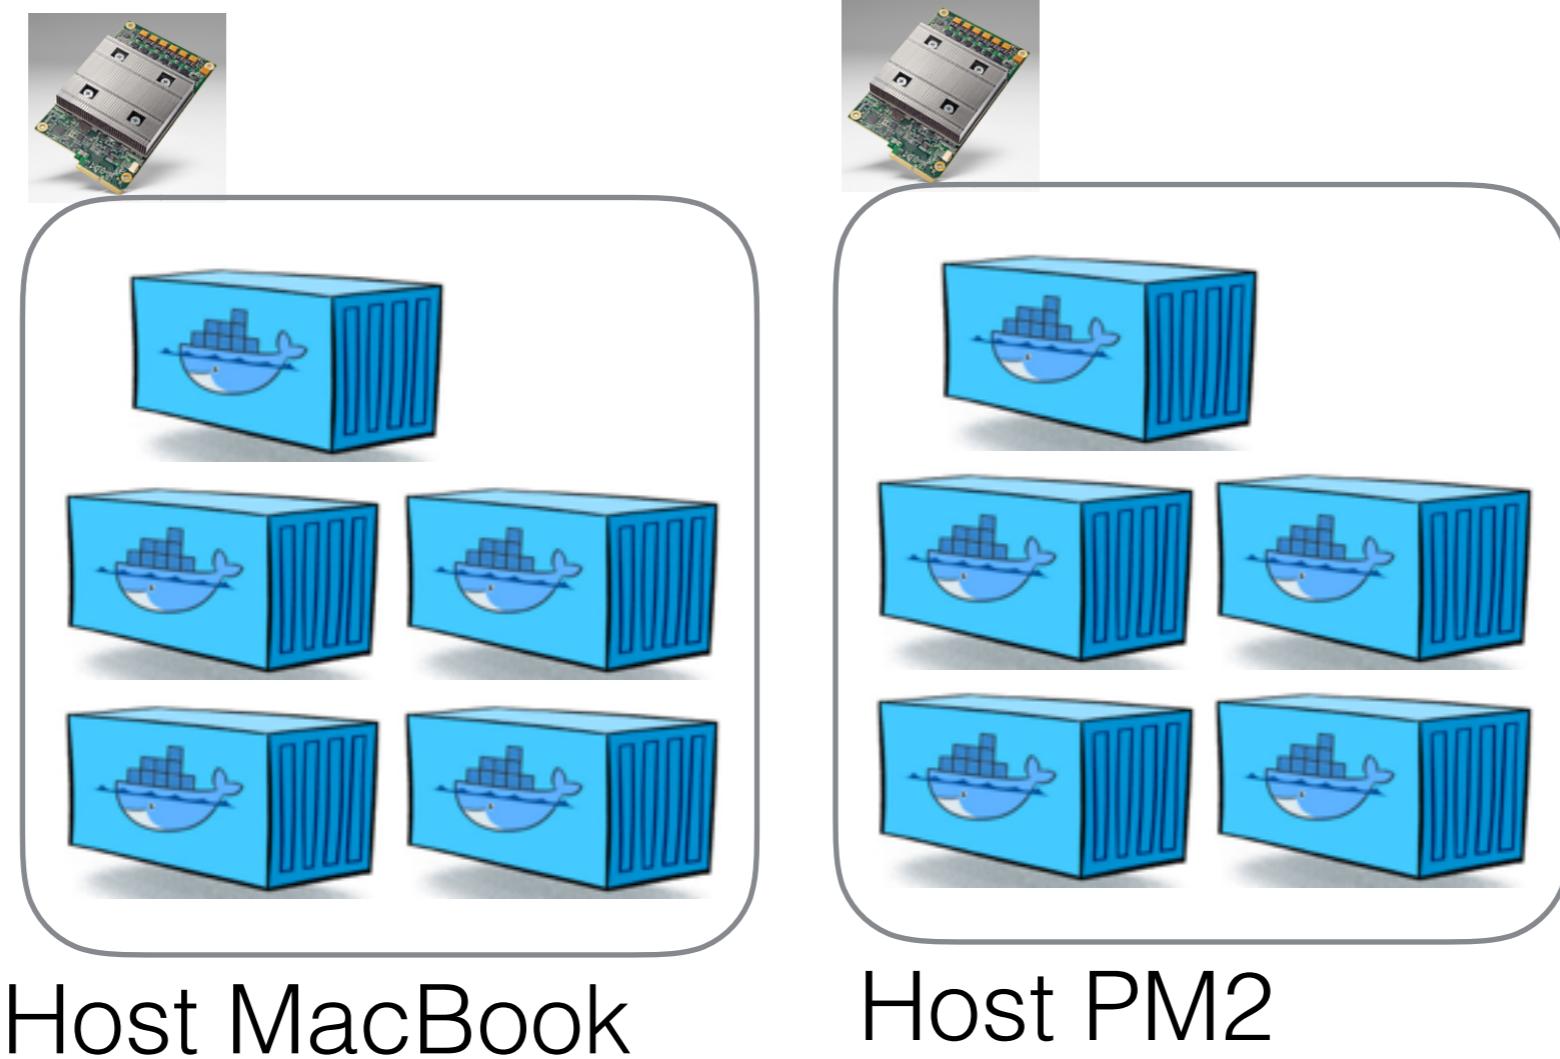
## GPU : “Open” issues

Assignee Sort

[Comment 0](#)[Comment 5](#)[Comment 2](#)[Comment 1](#)[Comment 5](#)[Comment 4](#)[Comment 0](#)[Comment 0](#)[Comment 10](#)[Comment 0](#)[Comment 2](#)[Comment 1](#)[Comment 1](#)

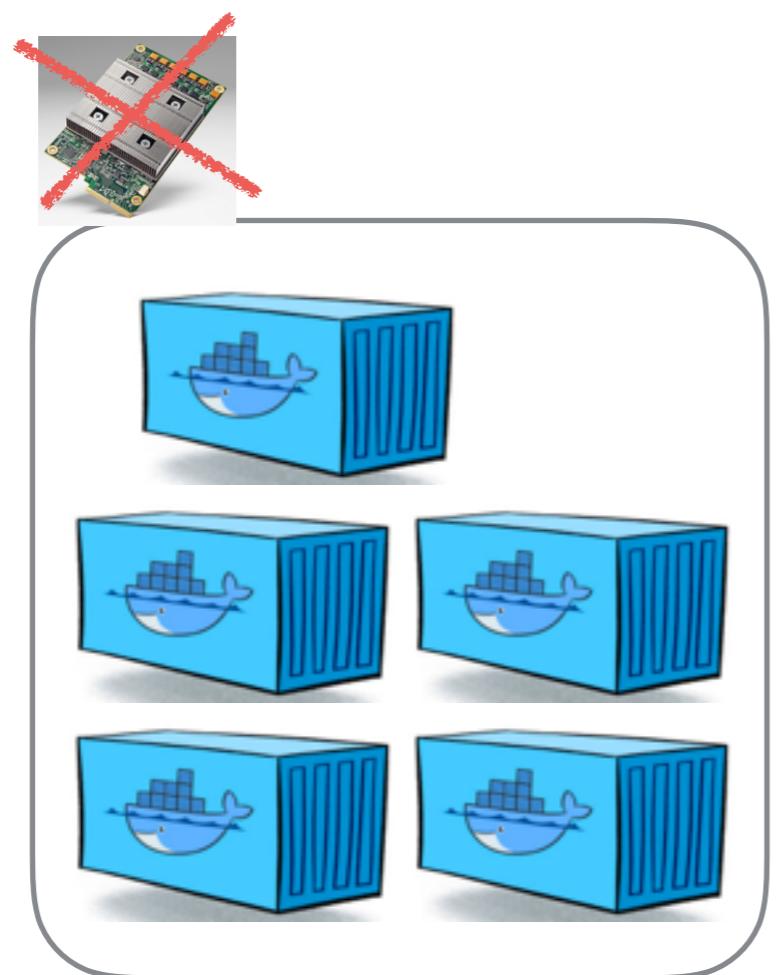
blog article :  
CUDA+Tensorflow+Docker : <https://goo.gl/dIXK2n>

# Demo for today...



My Experienced :  
FullSeg Player(MPEG-2) + GPU  
Media Player + H.264 + GPU  
=> Docker+GPU+Tensorflow ???

# Demo for today...

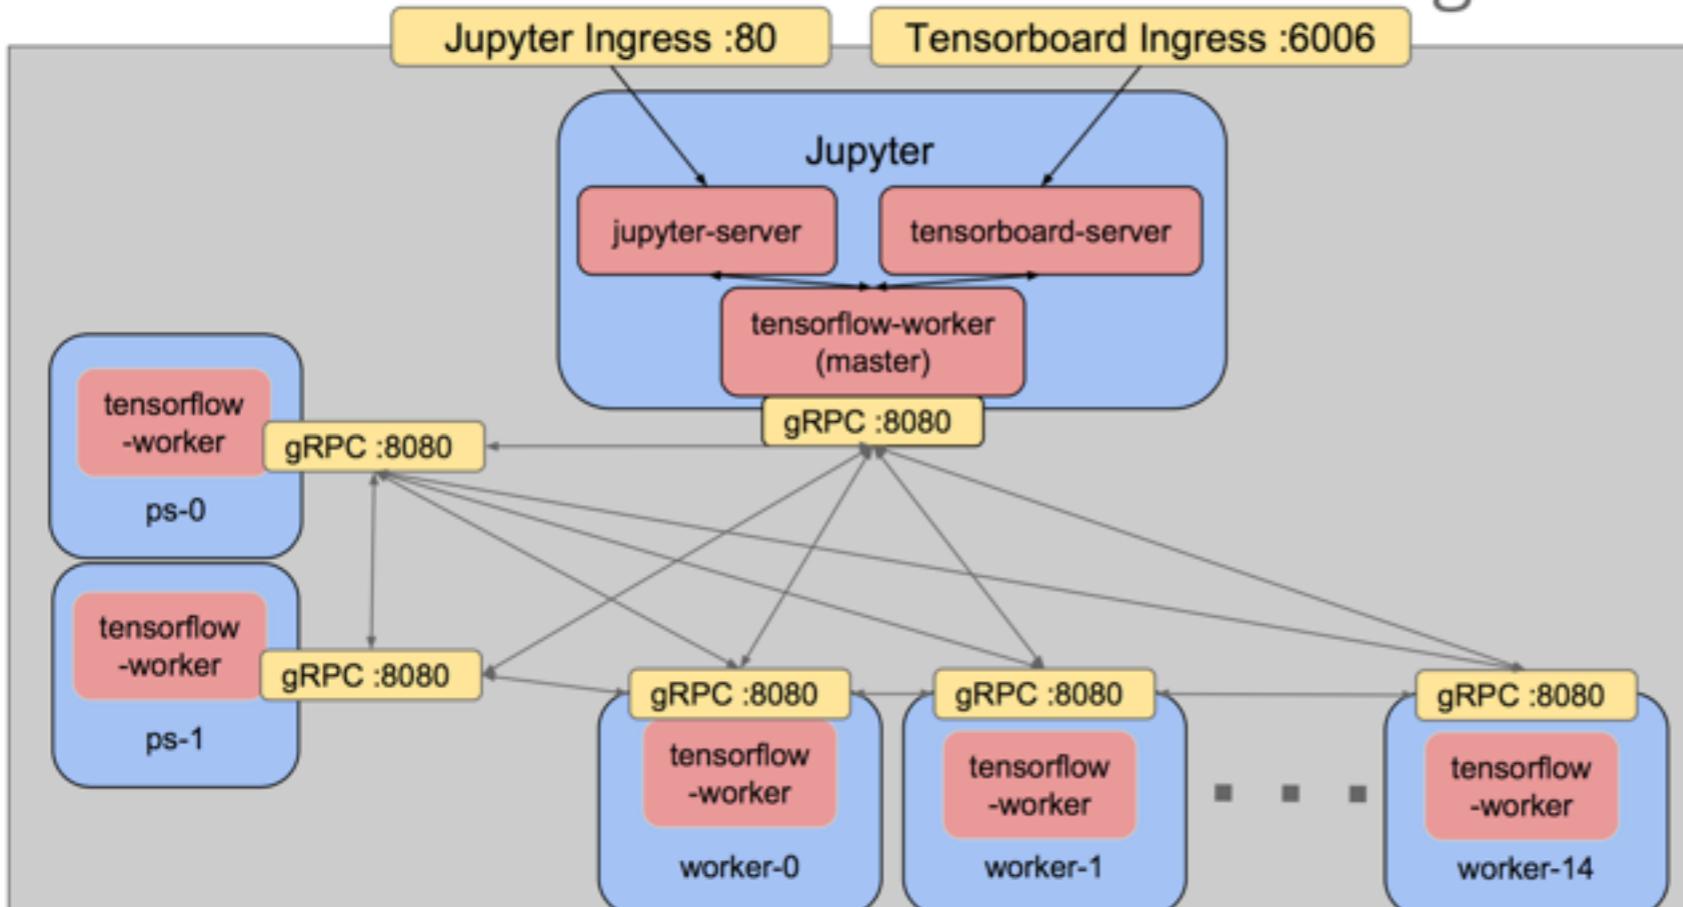


Host MacBook

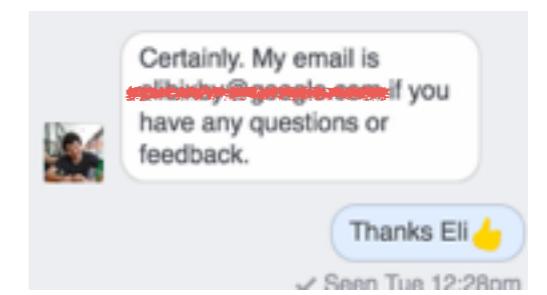
# TensorFlow Distributed Environment

: Basic

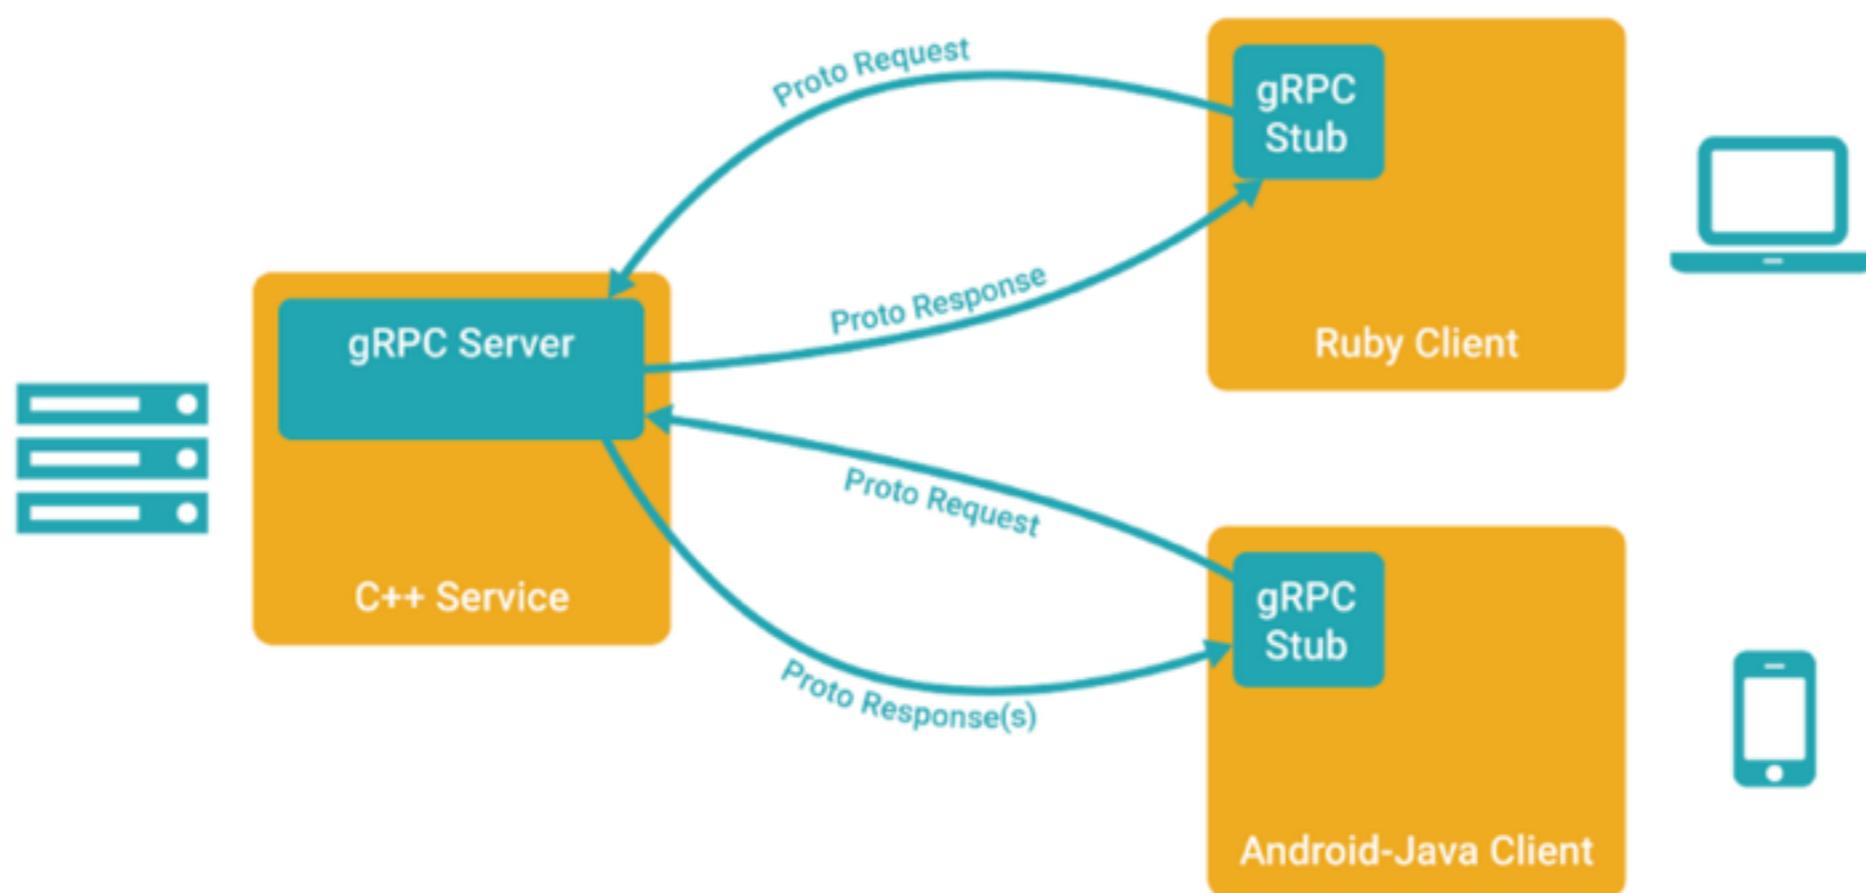
# Kubernetes as a Tensorflow Cluster Manager



[bit.ly/tensorflow-workshop](http://bit.ly/tensorflow-workshop)



# gRPC



# Hello distributed TensorFlow

```
# Start a TensorFlow server  
# as a single-process "cluster".  
  
$ python  
  
>>> import tensorflow as tf  
  
>>> c = tf.constant("Hello, distributed TensorFlow!")  
  
>>> server = tf.train.Server.create_local_server()  
  
Initialize HostPortsGrpcChannelCache for job local -> {localhost:55642}  
  
Started server with target: grpc://localhost:55642  
  
>>> sess = tf.Session(server.target) # Create a session on the server.  
  
>>> sess.run(c)  
'Hello, distributed TensorFlow!'
```

class **tf.train.Server**  
“ for use in distributed training.”

creates a single-process  
cluster, with an in-process  
server.

# Create a cluster

## cluster

a set of "tasks" that participate  
in the distributed execution

### tf.train.ClusterSpec construction

```
tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})
```

### Available tasks

/job:local/task:0  
/job:local/task:1

```
tf.train.ClusterSpec({  
    "worker": [  
        "worker0.example.com:2222",  
        "worker1.example.com:2222",  
        "worker2.example.com:2222"  
    ],  
    "ps": [  
        "ps0.example.com:2222",  
        "ps1.example.com:2222"  
    ]})
```

### task : worker0 TensorFlow server

...

/job:worker/task:0  
/job:worker/task:1  
/job:worker/task:2  
/job:ps/task:0  
/job:ps/task:1

# Create a tf.train.Server instance in each task

**Server(cluster, job\_name, task\_index)**

cluster:a set of "tasks" that participate in the distributed execution

```
# In task 0:
```

```
cluster = tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})  
server = tf.train.Server(cluster, job_name="local", task_index=0)
```

```
# In task 1:
```

```
cluster = tf.train.ClusterSpec({"local": ["localhost:2222", "localhost:2223"]})  
server = tf.train.Server(cluster, job_name="local", task_index=1)
```

cluster -> 1,000 server ?

# Specifying distributed devices in your model

```
with tf.device("/job:ps/task:0"):  
    weights_1 = tf.Variable(...)  
    biases_1 = tf.Variable(...)
```

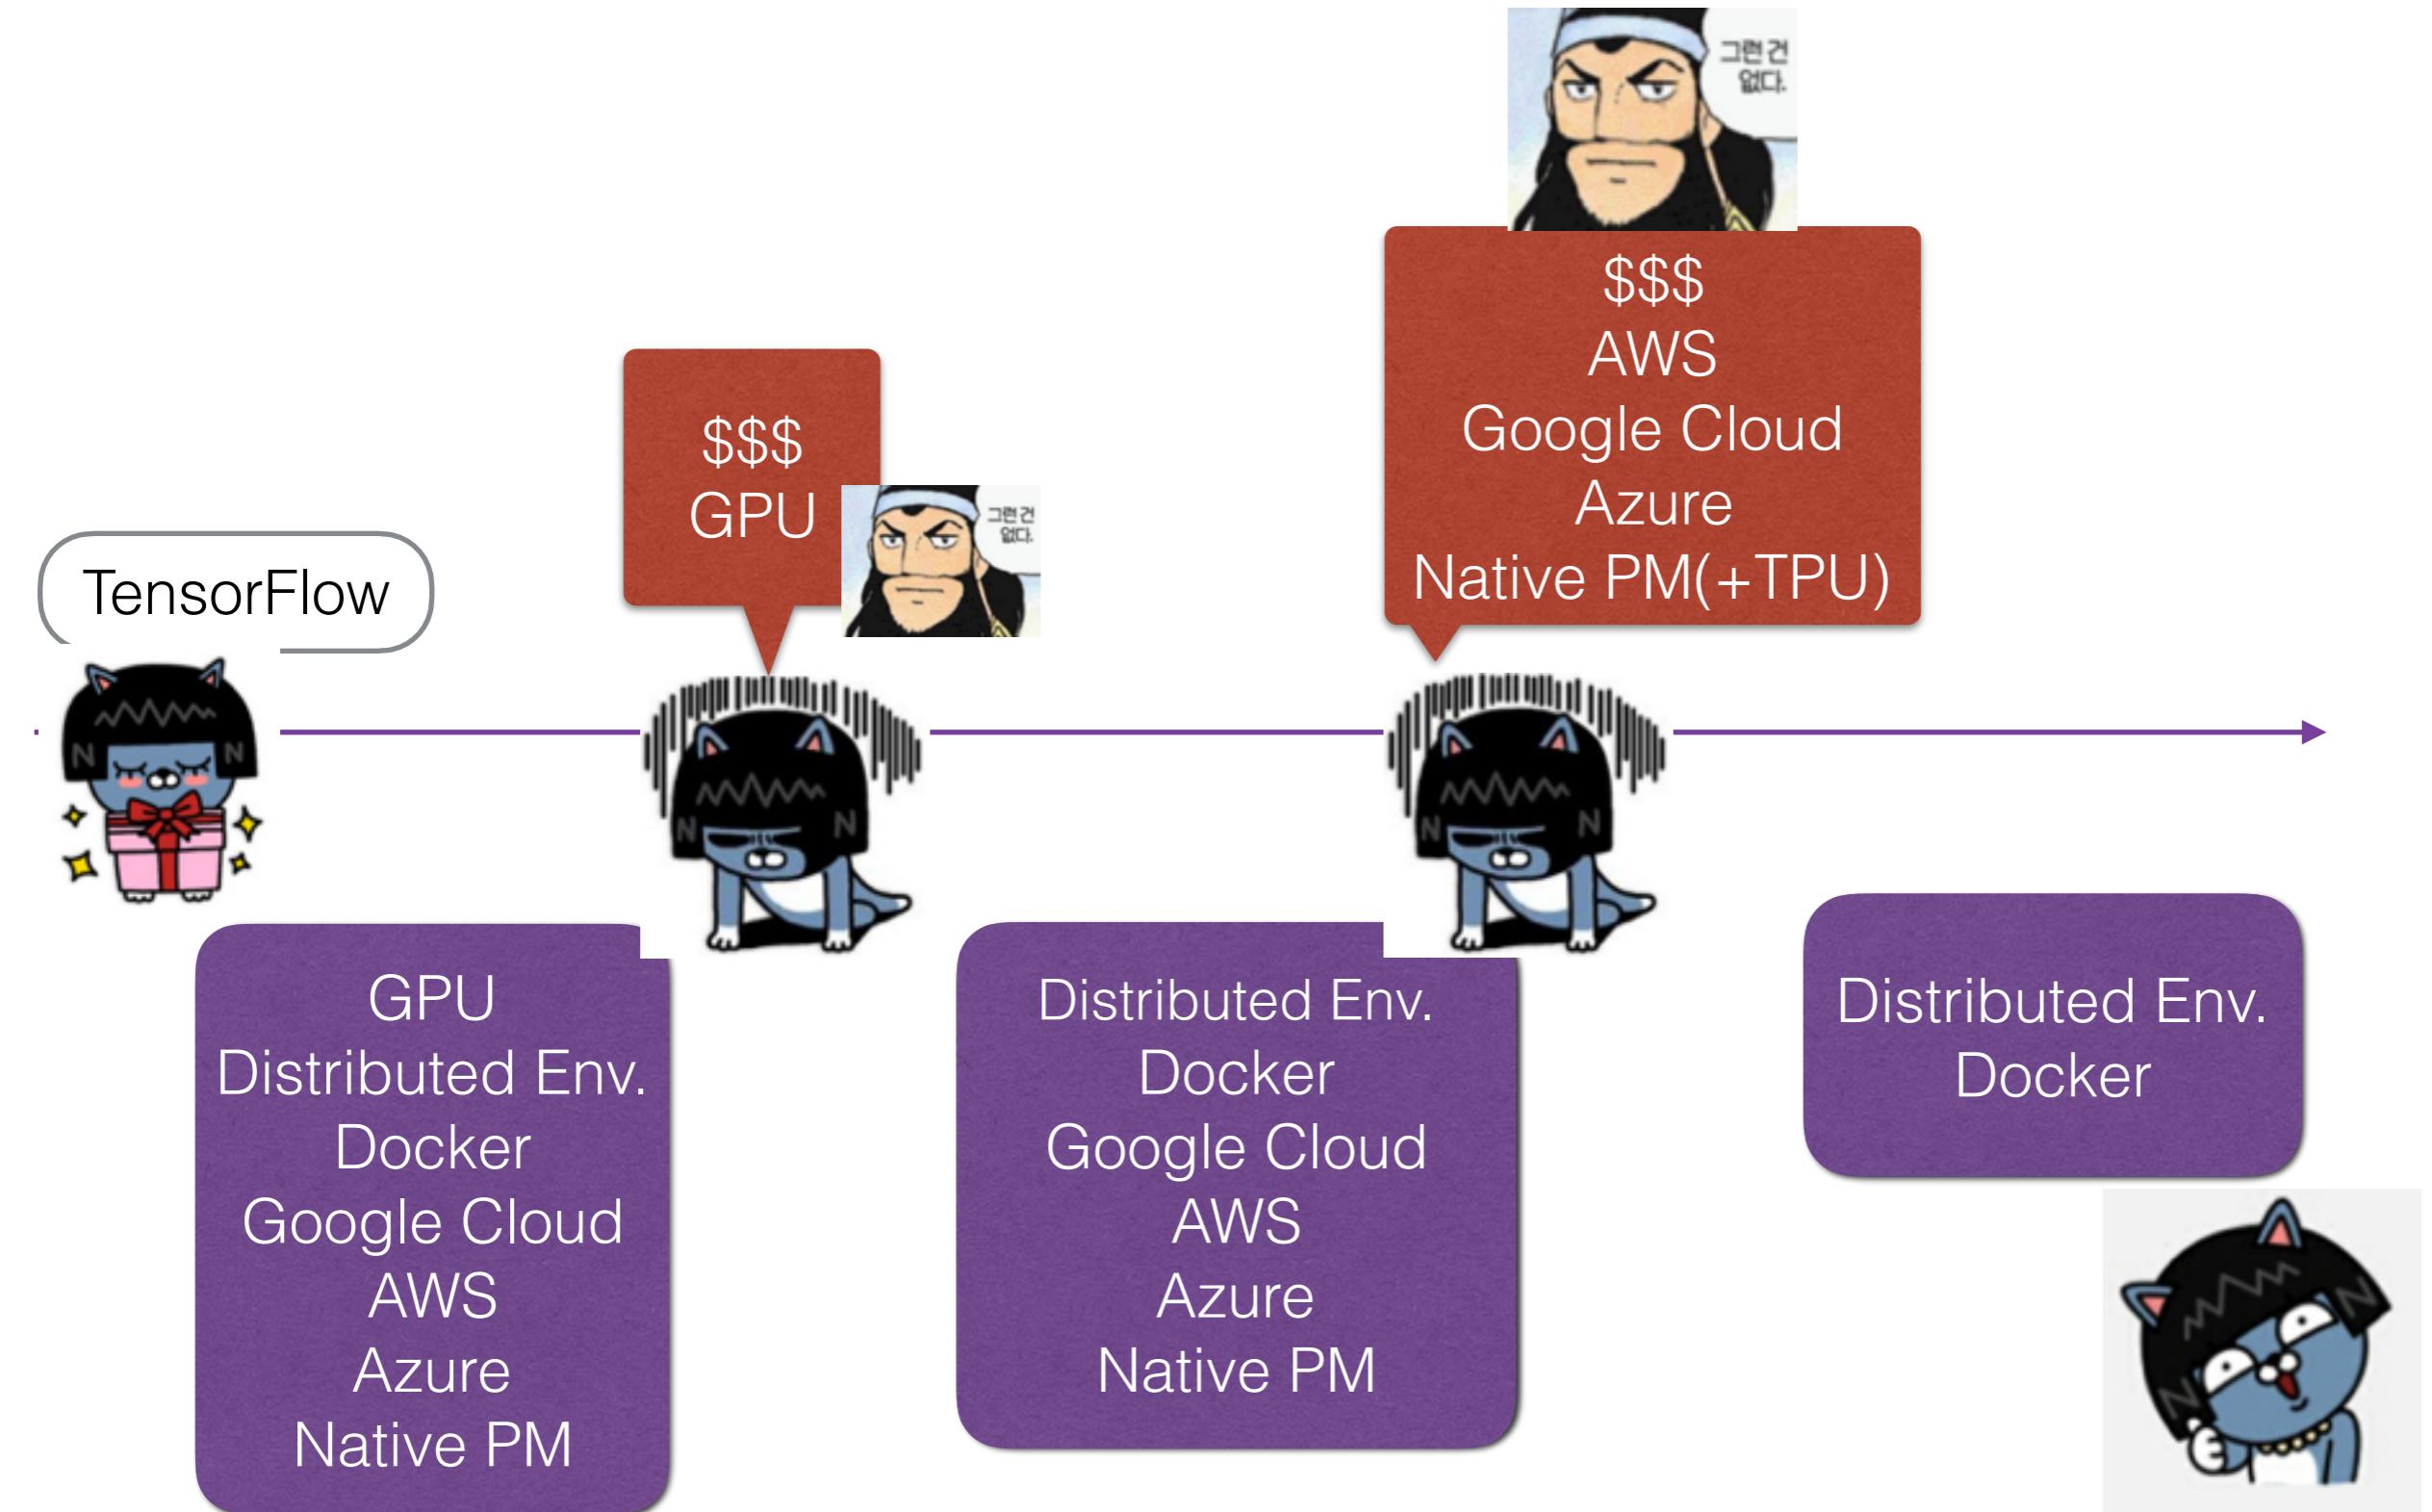
**tf.device() function that is used to specify whether ops run on the CPU or GPU.**

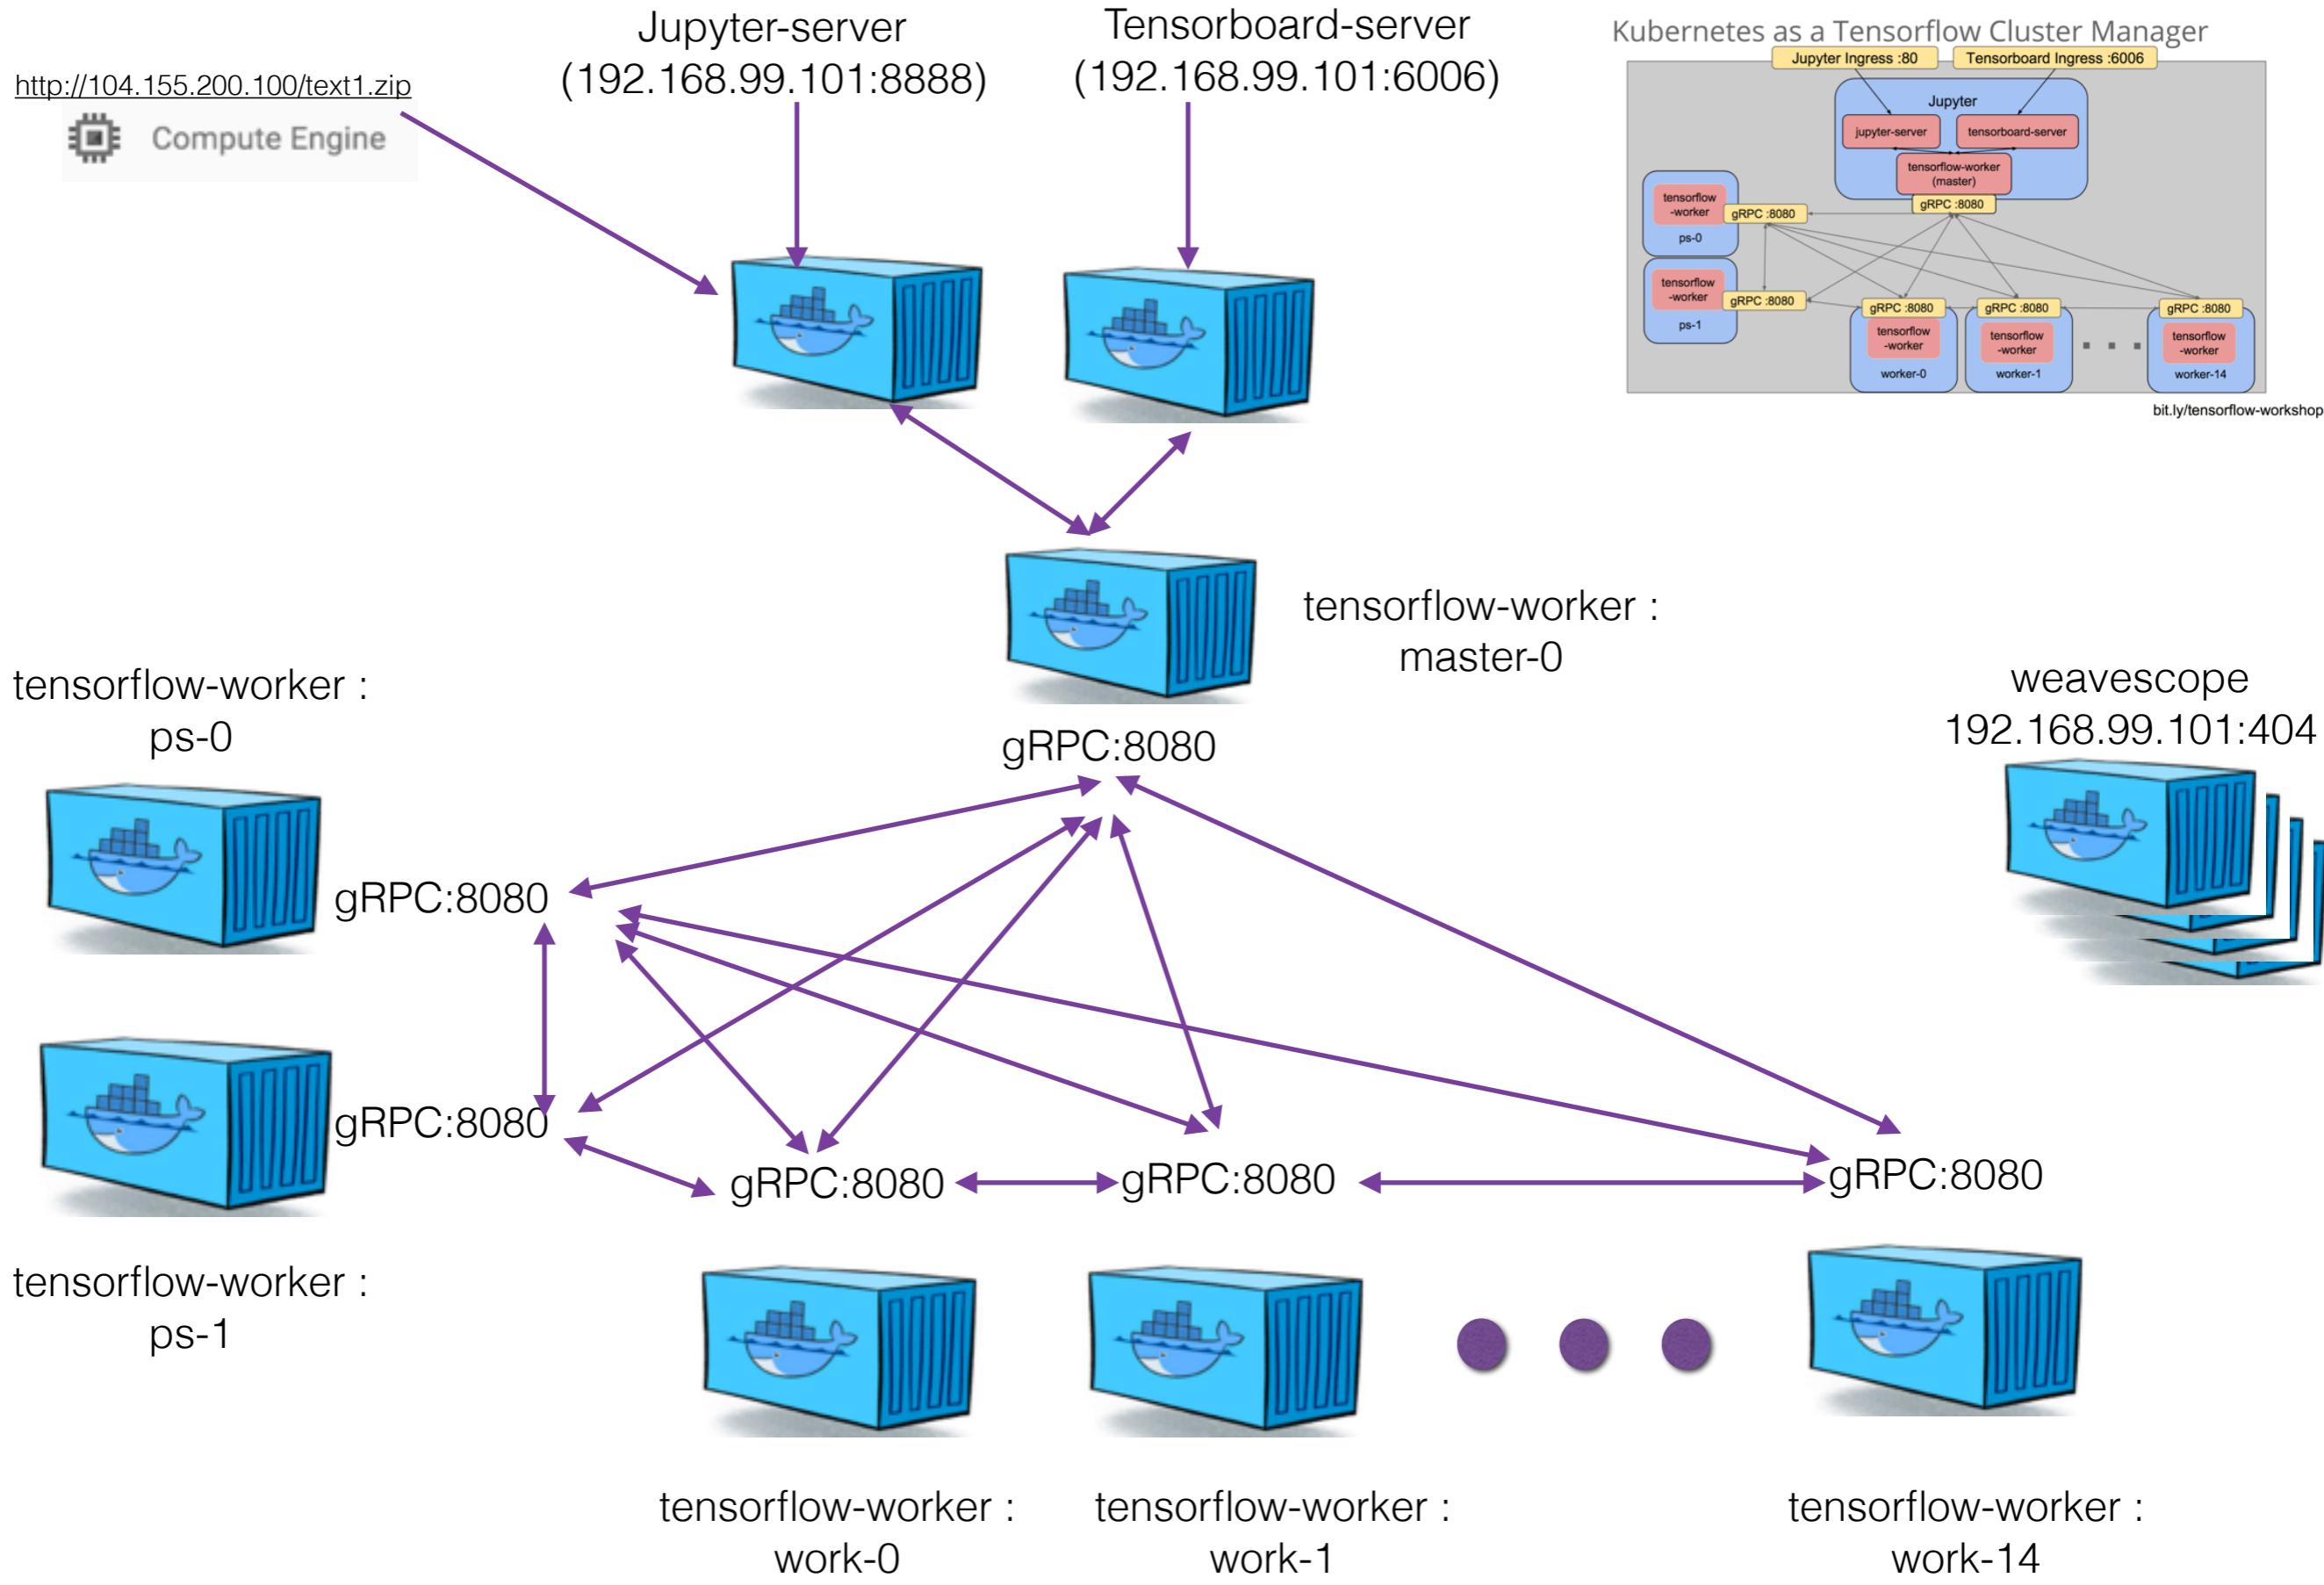
```
with tf.device("/job:worker/task:7"):  
    input, labels = ...  
    layer_1 = tf.nn.relu(tf.matmul(input, weights_1) + biases_1)  
    logits = tf.nn.relu(tf.matmul(layer_1, weights_2) + biases_2)  
    # ...  
    train_op = ...
```

**TensorFlow will insert the appropriate data transfers between the jobs(ps->worker, worker->ps)**

```
with tf.Session("grpc://worker7.example.com:2222") as sess:  
    for _ in range(10000):  
        sess.run(train_op)
```

Docker,  
Google Cloud Platform,  
AWS ..



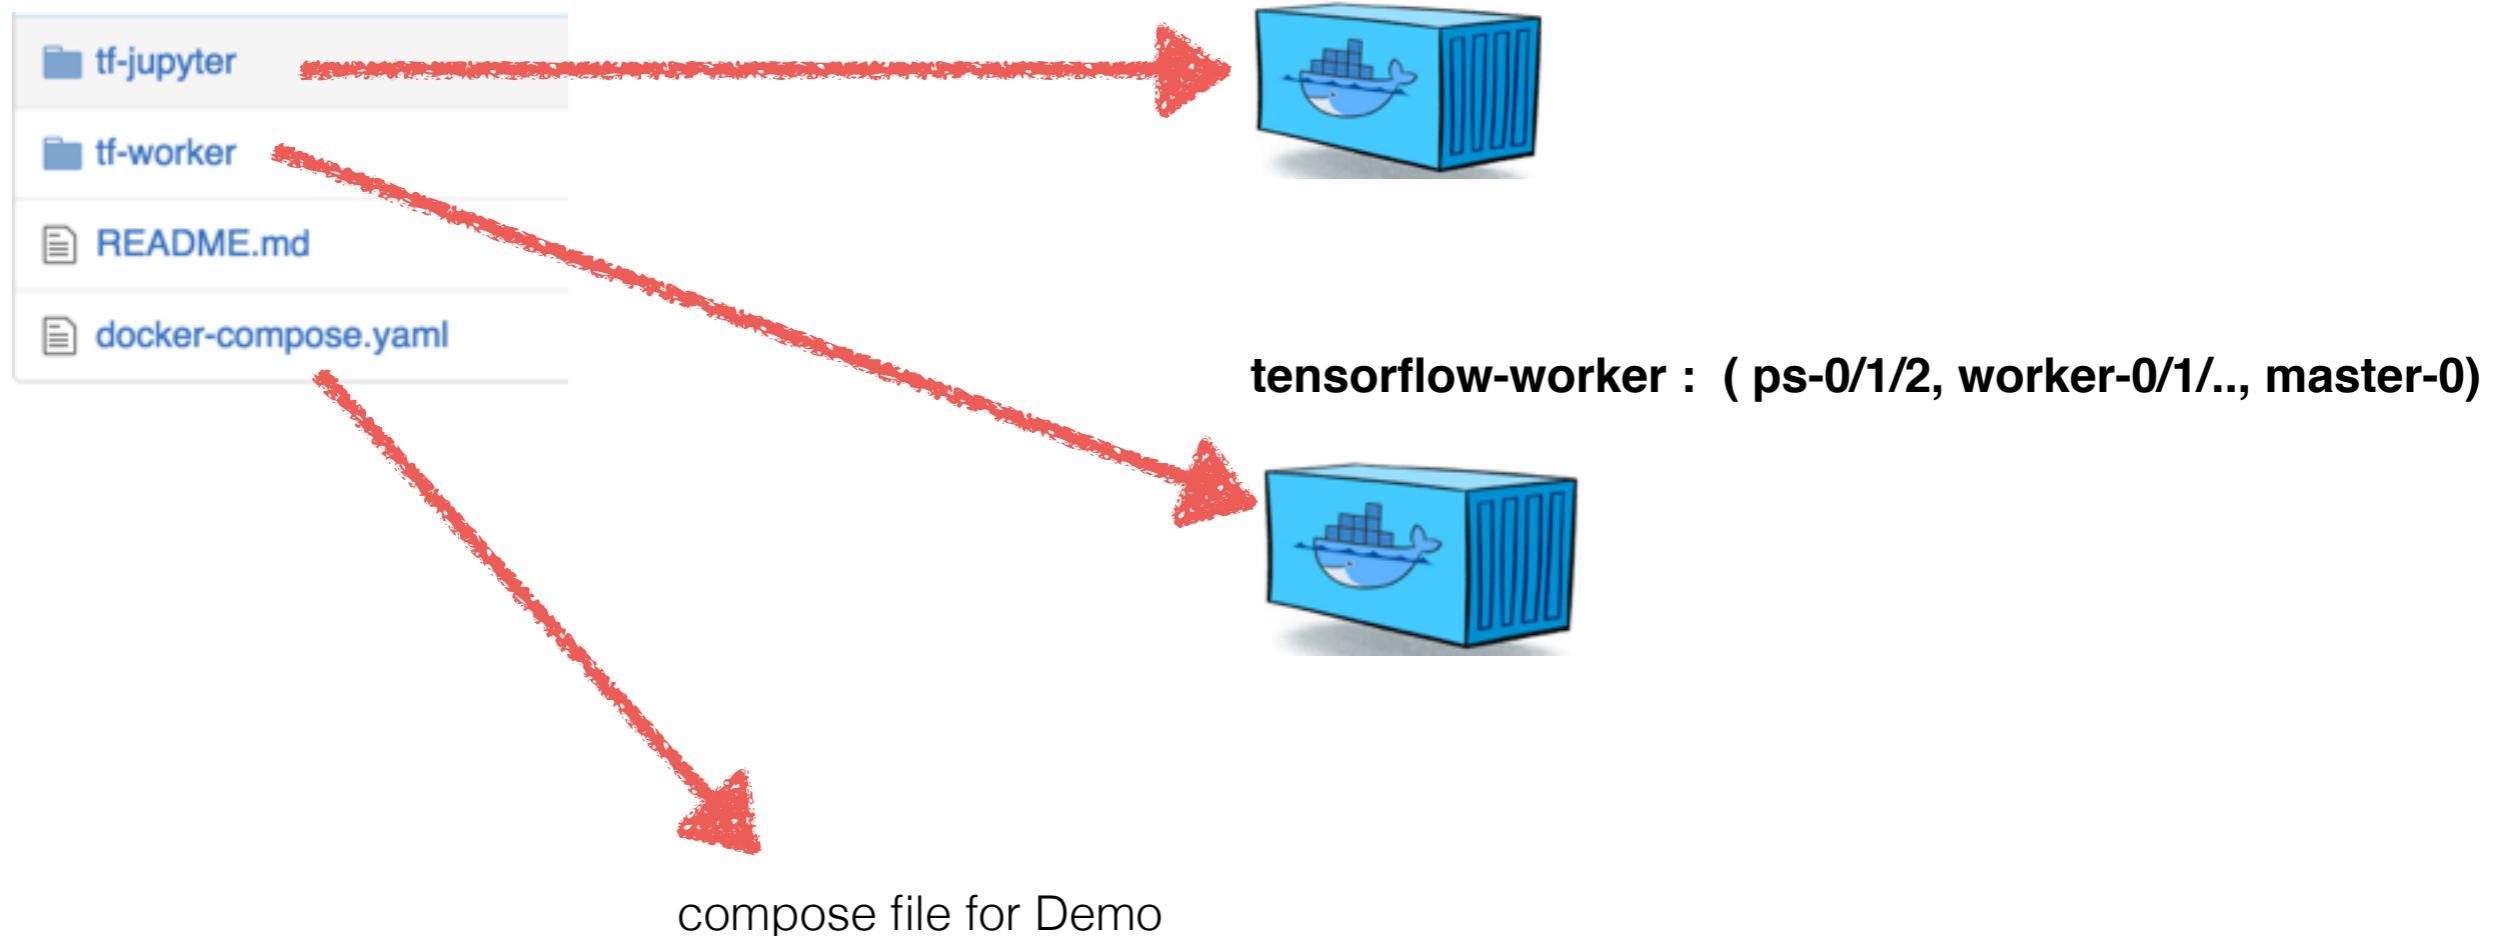


# Docker, Docker-compose...

Demo code :

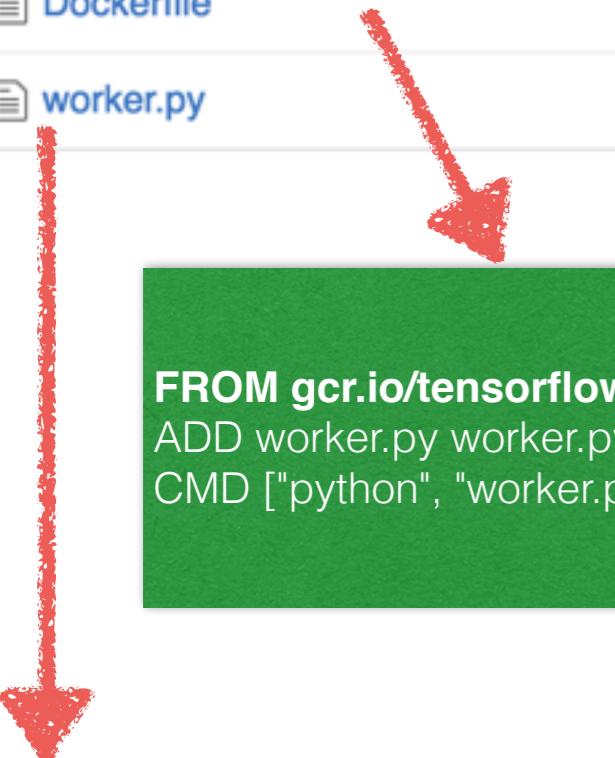
<https://github.com/bwahn/tensorflow-kr-docker>

github:



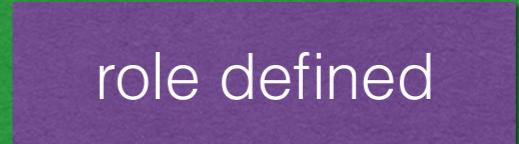
# tf-worker

```
..  
Dockerfile  
worker.py
```



```
FROM gcr.io/tensorflow/tensorflow:latest-devel  
ADD worker.py worker.py  
CMD ["python", "worker.py"]
```

```
def main(job_name, task_id):  
    server = tf.train.Server(  
        {  
            "ps": [  
                "ps-0:8080",  
                "ps-1:8080",  
                "ps-2:8080",  
                "ps-3:8080",  
            ],  
            "worker": [  
                "worker-0:8080",  
                "worker-1:8080",  
                "worker-2:8080",  
                "worker-3:8080",  
                "worker-4:8080",  
                "worker-5:8080",  
                "worker-6:8080",  
                "worker-7:8080",  
            ],  
            "master": [  
                "master-0:8080",  
            ],  
        },  
        job_name=job_name,  
        task_index=task_id  
    )  
    server.join()  
  
if __name__ == '__main__':  
    this_job_name, this_task_id = POD_NAME.split('-', 2)  
    main(this_job_name, int(this_task_id))
```



# ps-0 ( docker-compose.yaml )

```
version: '2'  
services:  
####  
ps-0:  
  build: ./tf-worker/  
  container_name: ps-0  
image: ps-0:0.1  
  ports:  
    - "8080"  
  environment:  
    POD_NAME: ps-0  
....
```

gRPC port

POD\_NAME env in worker.py

```
$ docker-compose -f docker-compose.yaml build ps-0
```

```
Building ps-0
```

```
Step 1 : FROM gcr.io/tensorflow/tensorflow:latest-devel
```

```
---> c3efccc5f94f
```

```
Step 2 : ADD worker.py worker.py
```

```
---> Using cache
```

```
---> f9c2d840b051
```

```
Step 3 : CMD python worker.py
```

```
---> Using cache
```

```
---> 57a6c024580c
```

```
Successfully built 57a6c024580c
```

```
$ docker images | grep ps -0
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ps-0	0.1	57a6c024580c	2 days ago	2.38 GB

```
version: '2'  
services:  
####  
ps-0:  
build: ./tf-worker/  
container_name: ps-0  
image: ps-0:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: ps-0
```

```
ps-1:  
build: ./tf-worker/  
container_name: ps-1  
image: ps-1:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: ps-1
```

```
ps-2:  
build: ./tf-worker/  
container_name: ps-2  
image: ps-2:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: ps-2
```

```
ps-3:  
build: ./tf-worker/  
container_name: ps-3  
image: ps-3:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: ps-3
```

```
####  
worker-0:  
build: ./tf-worker/  
container_name: worker-0  
image: worker-0:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-0
```

```
worker-1:  
build: ./tf-worker/  
container_name: worker-1  
image: worker-1:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-1
```

```
worker-2:  
build: ./tf-worker/  
container_name: worker-2  
image: worker-2:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-2
```

```
worker-3:  
build: ./tf-worker/  
container_name: worker-3  
image: worker-3:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-3
```

```
worker-4:  
build: ./tf-worker/  
container_name: worker-4  
image: worker-4:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-4
```

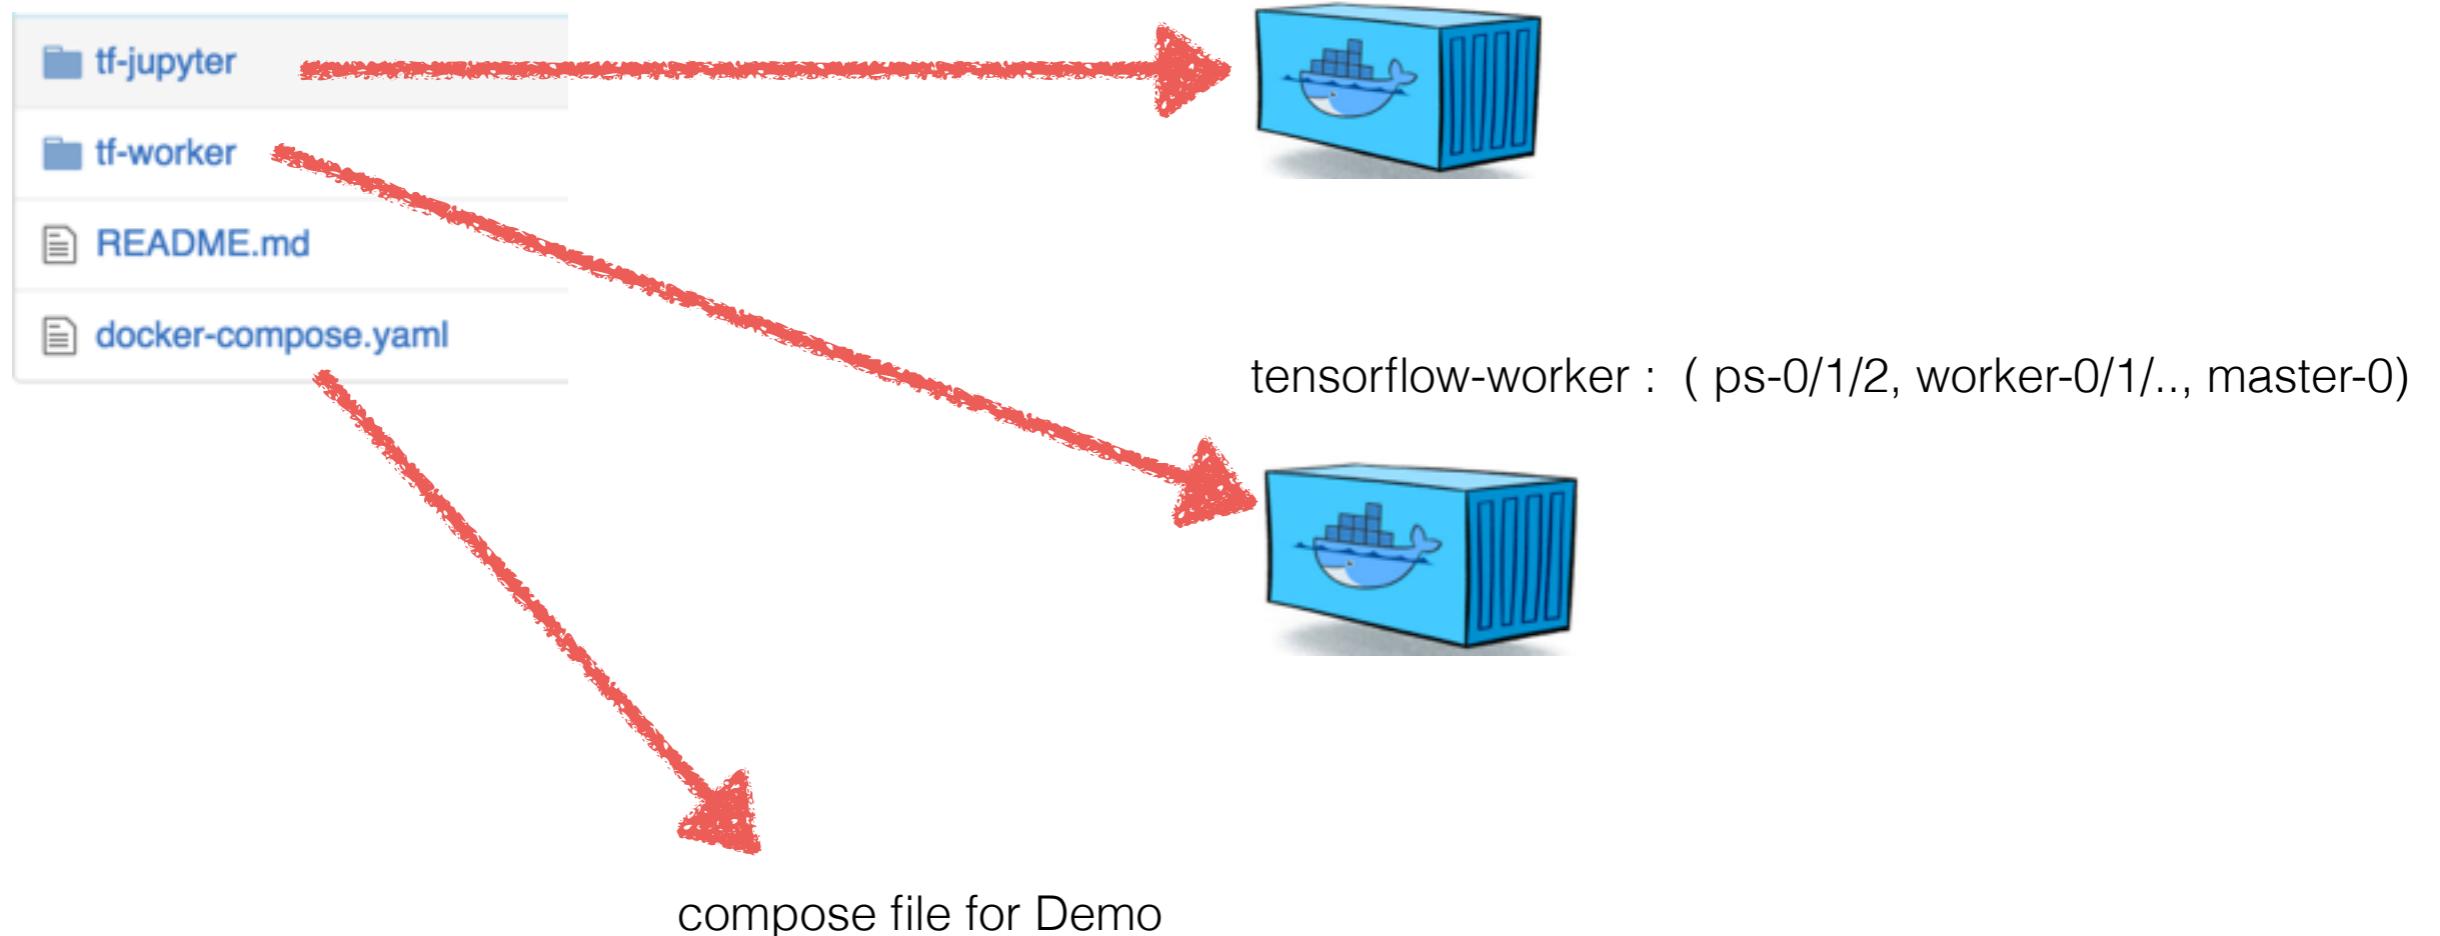
```
worker-5:  
build: ./tf-worker/  
container_name: worker-5  
image: worker-5:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-5
```

```
worker-6:  
build: ./tf-worker/  
container_name: worker-6  
image: worker-6:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-6
```

```
worker-7:  
build: ./tf-worker/  
container_name: worker-7  
image: worker-7:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: worker-7
```

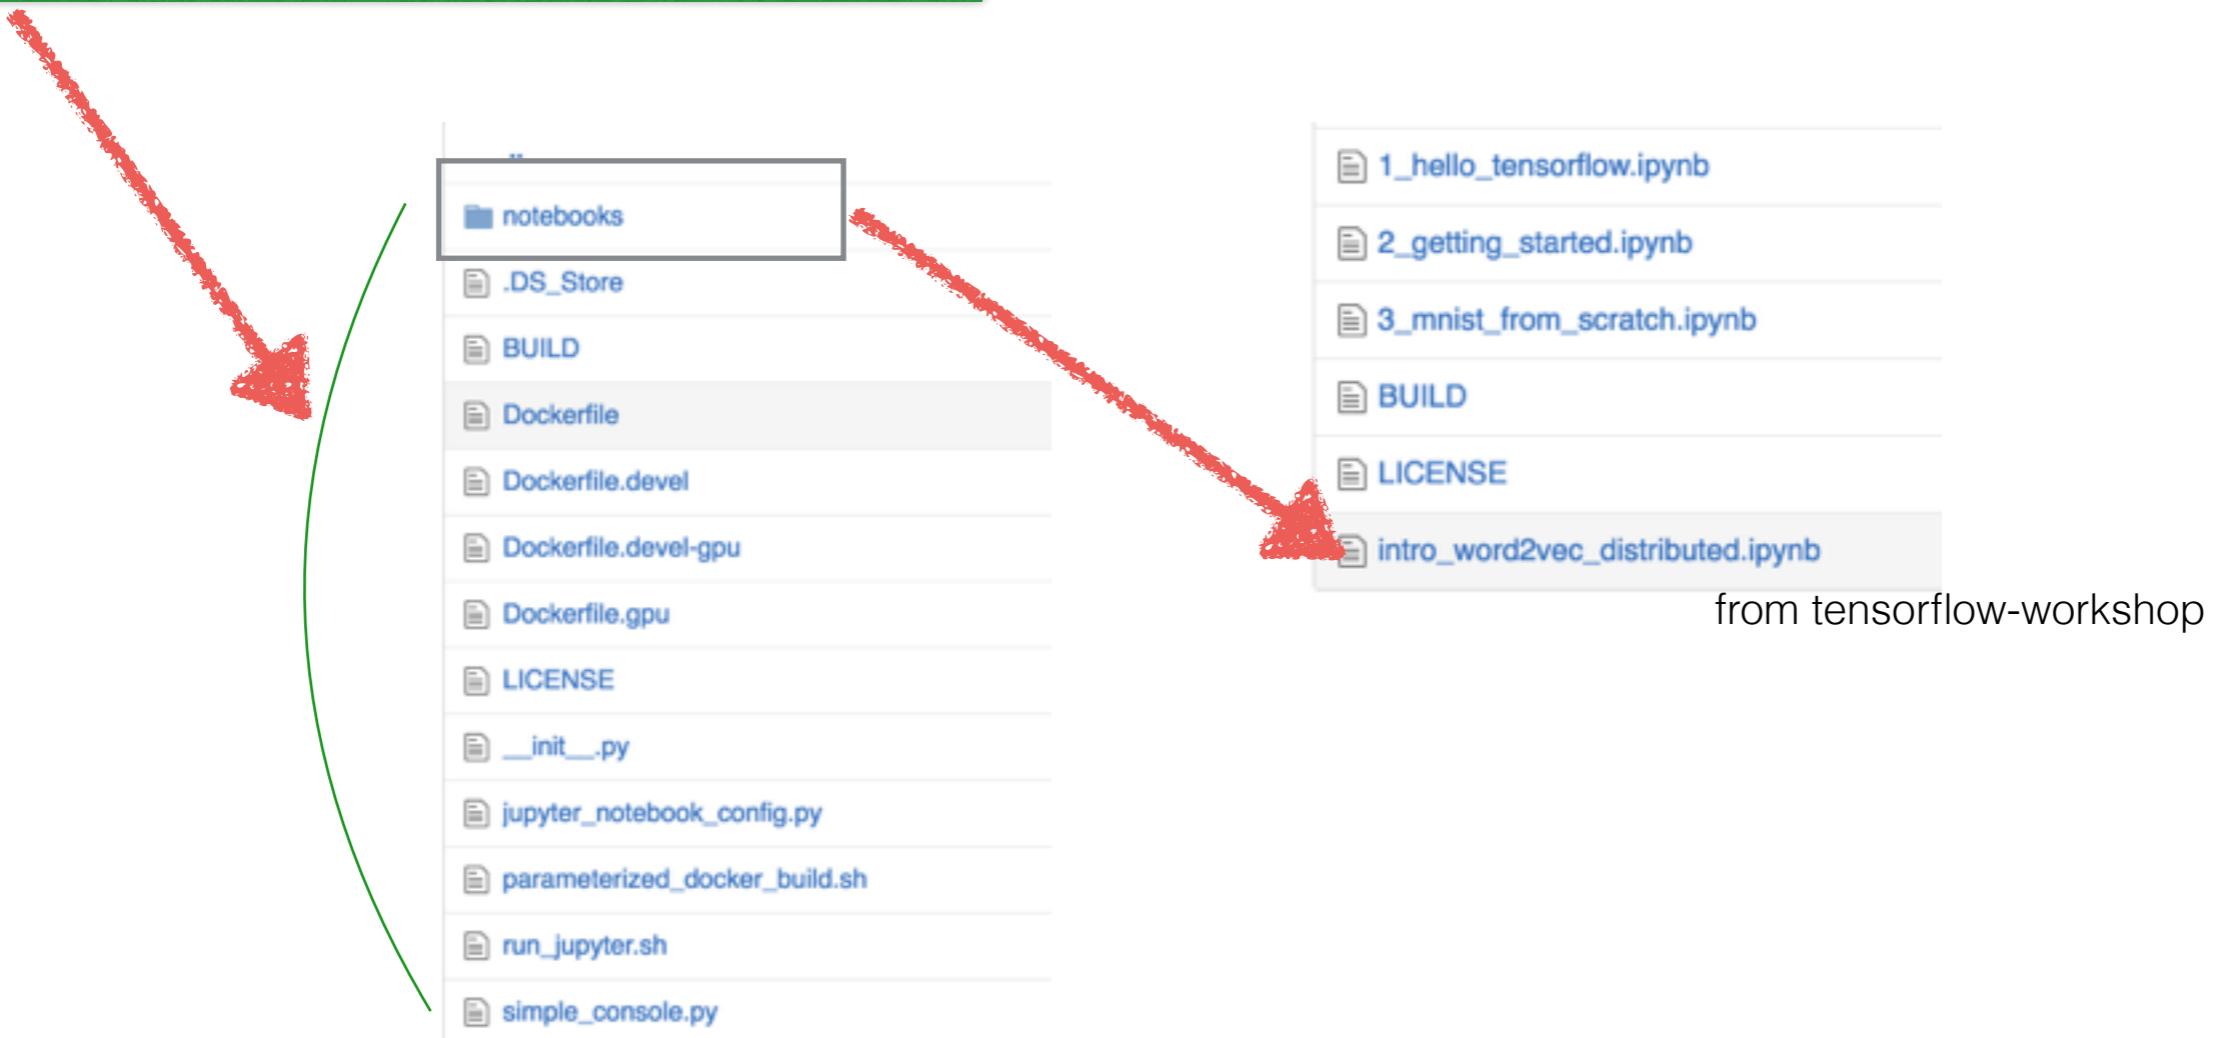
```
####  
master-0:  
build: ./tf-worker/  
container_name: master-0  
image: master-0:0.1  
ports:  
- "8080"  
environment:  
POD_NAME: master-0
```

github:



# tf-jupyter

```
// TensorFlow github  
$ git clone https://github.com/tensorflow/tensorflow.git  
$ cd /tensorflow-git/tensorflow/tools/docker
```



```
$ docker-compose -f docker-compose.yaml build jupyter-server
Building jupyter-server
Step 1 : FROM ubuntu:14.04
--> 8f1bd21bd25c
Step 2 : MAINTAINER Craig Citro <craigcitro@google.com>
--> Using cache
--> c87f53e1fc9
Step 3 : RUN apt-get update && apt-get install -y curl libfreetype6-dev libpng12-dev python-scipy && apt-get clean && rm -rf /var/lib/apt/lists/*
--> Using cache
--> a5abc6a5ed5a
Step 4 : RUN curl -O https://bootstrap.pypa.io/get-pip.py && python get-pip.py && rm get-pip.py
--> Using cache
--> 39df3de28486
Step 5 : RUN pip --no-cache-dir install ipykernel jupyter matplotlib && python
--> Using cache
--> ce0e0004e0a6
Step 6 : ENV TENSORFLOW_VERSION 0.8.0
--> Using cache
--> 80791608c082
Step 7 : RUN pip --no-cache-dir install http://storage.googleapis.com/tensorflow/linux/cpu/tensorflow
--> Using cache
--> 980ebdfed88c
Step 8 : COPY jupyter_notebook_config.py /root/.jupyter/
--> Using cache
--> 2e2aa264d165
Step 9 : COPY notebooks /notebooks
--> Using cache
--> 3b6409cc98a9
Step 10 : COPY run_jupyter.sh /
--> Using cache
--> b6cfb46577ed
Step 11 : EXPOSE 6006
--> Using cache
--> e4010cf837c6
Step 12 : EXPOSE 8888
--> Using cache
--> 0266ba056034
Step 13 : WORKDIR "/notebooks"
--> Using cache
--> 48983eb7b7af
Step 14 : CMD /run_jupyter.sh
--> Using cache
--> 364efc0d61ce
Successfully built 364efc0d61ce
```

```
$ docker images | grep tf-jupyter-server
REPOSITORY      TAG        IMAGE ID      CREATED       SIZE
tf-jupyter-server    0.1        364efc0d61ce   4 days ago   722 MB
```

//docker-compose.yaml

### jupyter-server:

build:

  context: ./tf-jupyter/

  container\_name: tf-jupyter-server

  image: tf-jupyter-server:0.1

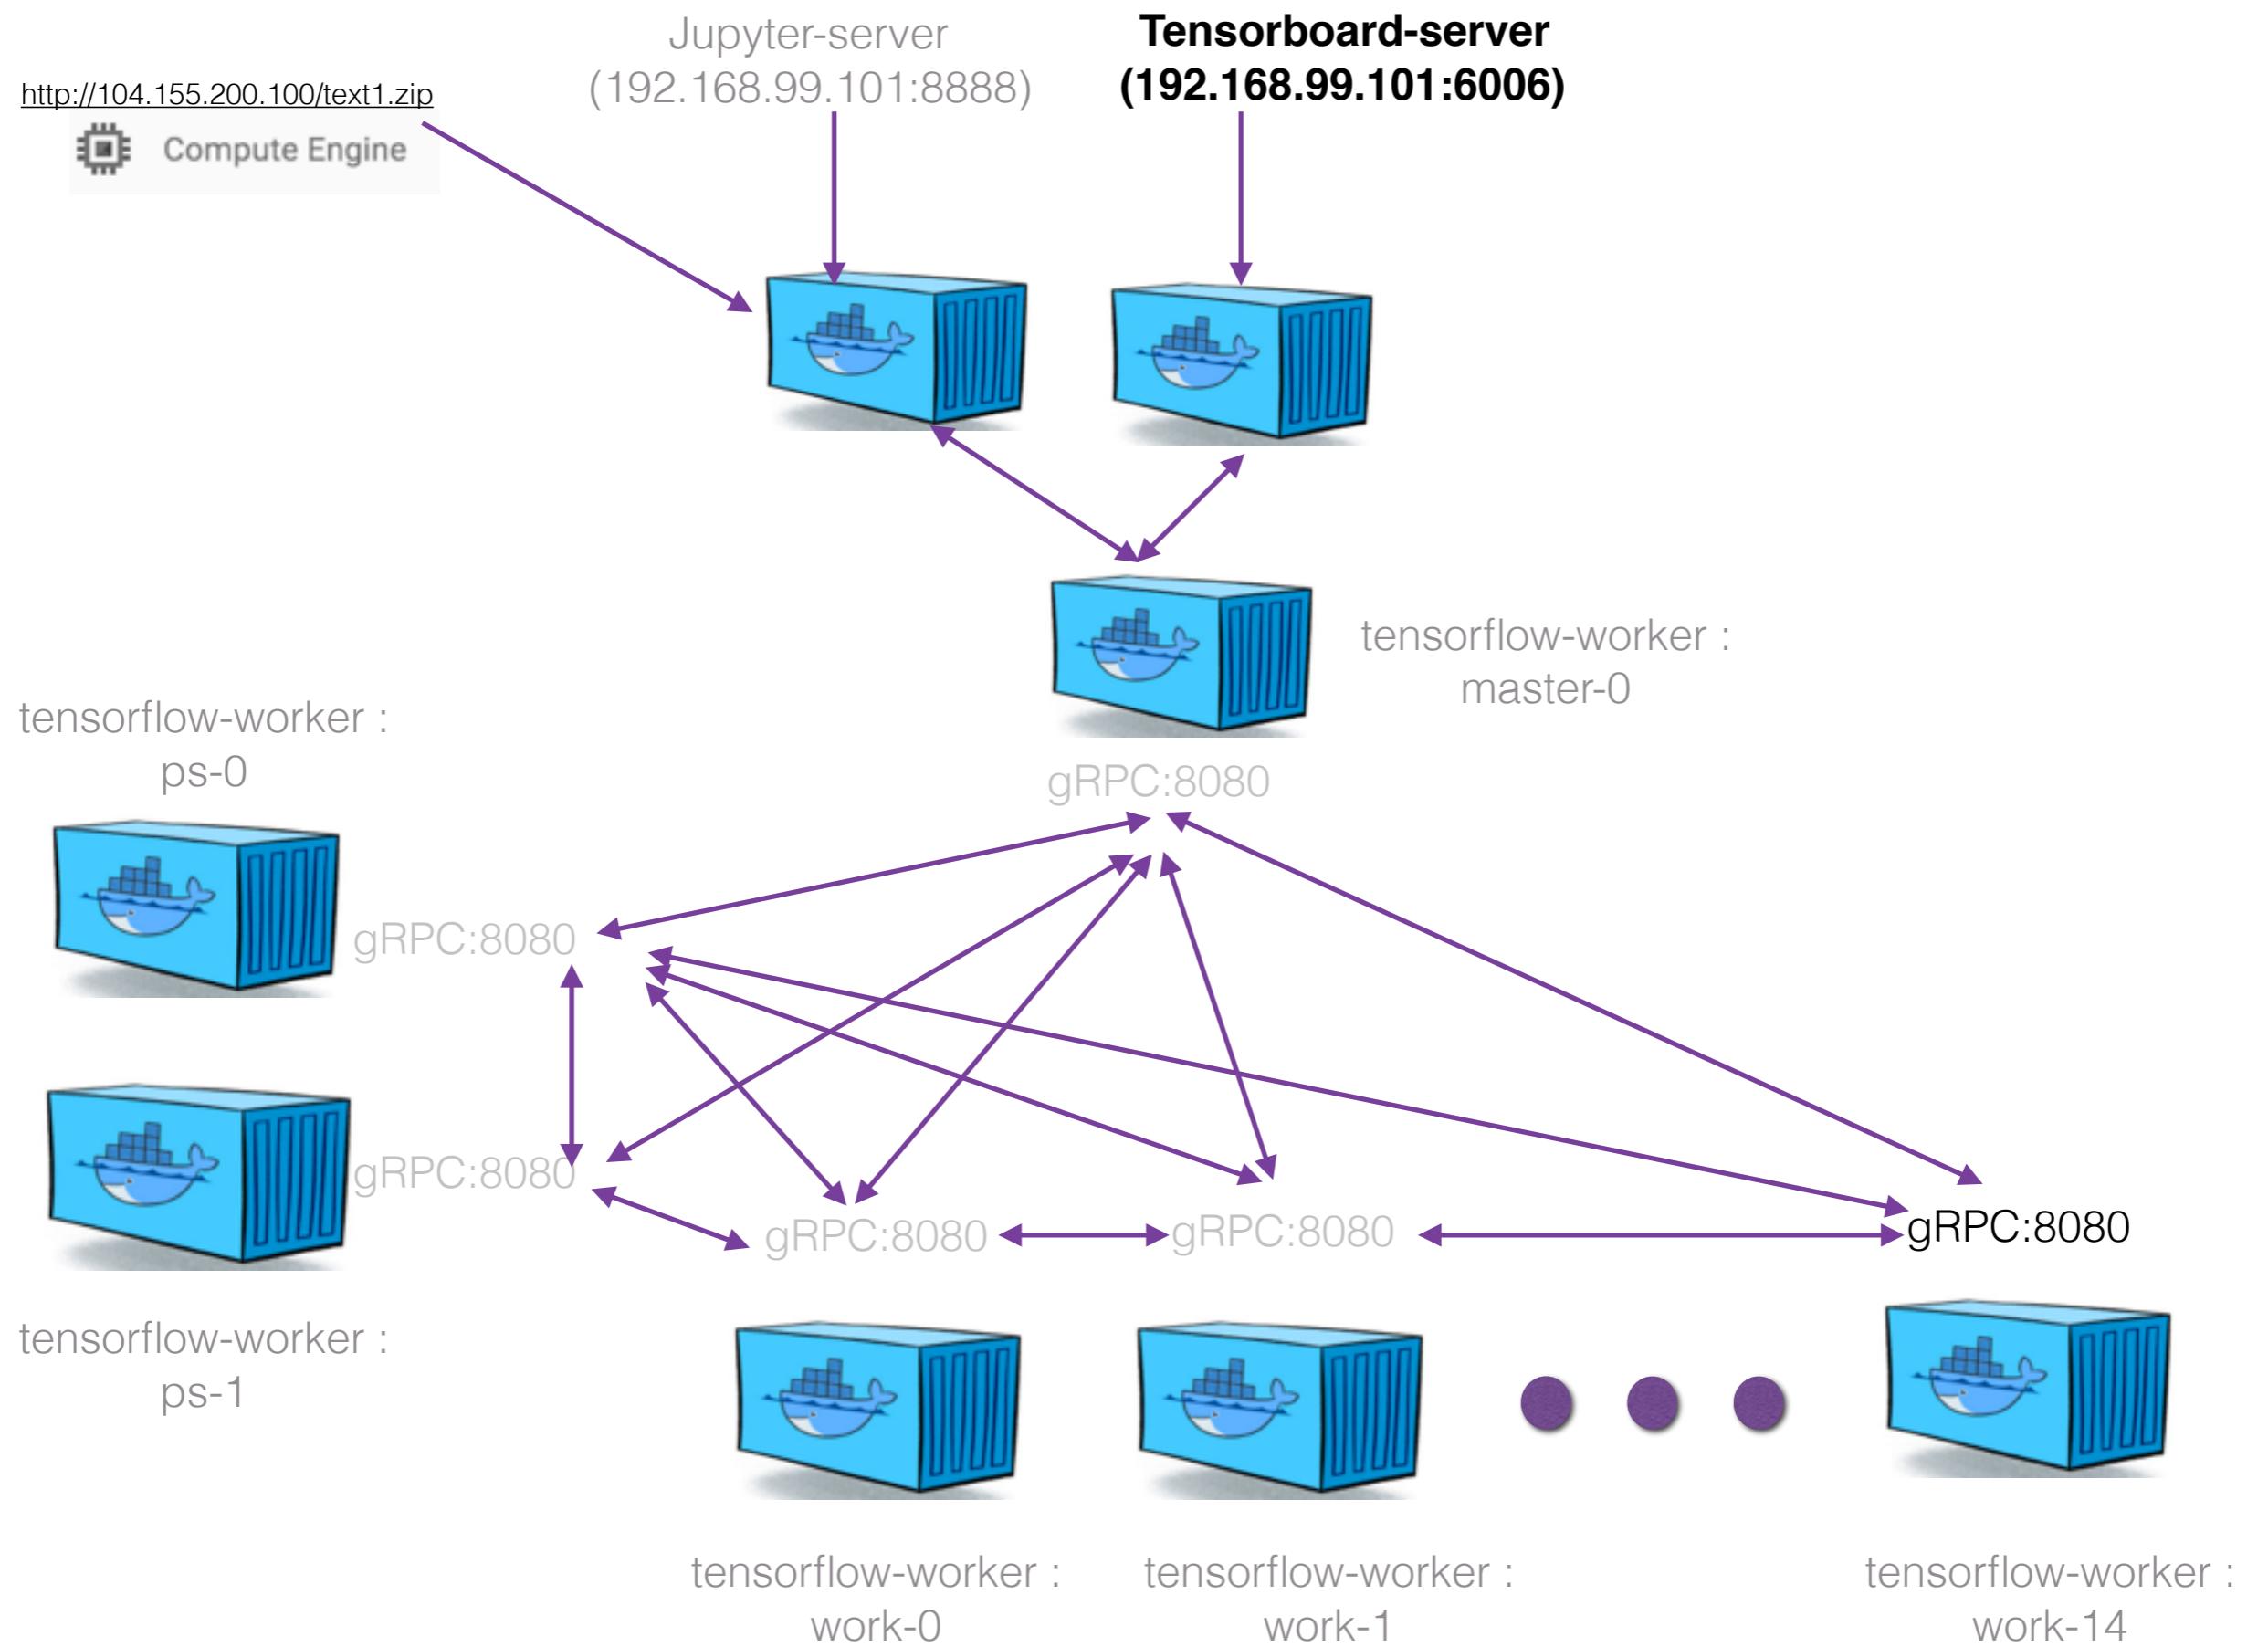
  ports:

- "8888:8888"

- "80:80"

  volumes:

- /tmp/tf/tensorflow-logs:/var/log/tensorflow



# Tensorboard-server

```
//docker-compose.yaml
...
tf-tensorboard-server:
  image: gcr.io/tensorflow/tensorflow:latest-devel
  container_name: tf-tensorboard-server
  ports:
    - "6006:6006"
  command: /tensorflow/bazel-bin/tensorflow/tensorboard/tensorboard --logdir=/var/log/tensorflow
  volumes:
    - /tmp/tf/tensorflow-logs:/var/log/tensorflow
```

## Summary

- Tensorflow Distributed Environment ( with Docker )

## Future

- Native Clustering (several hosts ): Docker Swarm, RancherOS
- Docker + GPU Architecture & Test Environment

## Q&A